



Digital Egypt Pioneers Initiative (DEPI) Program

Microsoft Machine Learning Engineering Track

CAI4_AIS2_G2

Customer Support RAG-Powered Intelligent Chatbot

Final Project Documentation

Prepared by:

Ahmed Farahat Ibrahim

Bassel Wael Kmal

Mahmoud Osama Ahmed

Mohamed Abdelfattah Mohamed

Sarah Abdulhamid Abdulkhalik

Nora Mohamed Mohamed

Submission Date: 25/05/2026

Table of Contents:

- **Abstract**
- **Chapter 1: Introduction**
 - 1.1 Project Overview
 - 1.2 Problem Statement
 - 1.3 Project Objectives
 - 1.4 Scope of the Project
 - 1.5 Technologies Used
- **Chapter 2: Data Collection and Preprocessing**
 - 2.1 Data Sources
 - 2.2 Knowledge Base Preparation
 - 2.3 Data Ingestion Pipeline
 - 2.4 Text Cleaning and Arabic Normalization
 - 2.5 Tokenization and Chunking
 - 2.6 Exploratory Data Analysis Summary
 - 2.7 Final Processed Corpus
- **Chapter 3: System Architecture and RAG Pipeline**
 - 3.1 Overall System Architecture
 - 3.2 Retrieval-Augmented Generation Architecture
 - 3.3 Embedding Generation Workflow
 - 3.4 Vector Store and Semantic Retrieval
 - 3.5 Query Optimization
 - 3.6 Grounded Response Generation
 - 3.7 Conversational Memory
 - 3.8 API Design and Frontend Integration
- **Chapter 4: Model Development and Optimization**
 - 4.1 Embedding Model Selection
 - 4.2 Vector Database Configuration
 - 4.3 Retriever Configuration
 - 4.4 Prompt Engineering Strategy
 - 4.5 Azure OpenAI Integration
 - 4.6 PSO-Based Retrieval Optimization
 - 4.7 Response Generation Process

- **Chapter 5: Evaluation and Testing**

- 5.1 Evaluation Methodology

- 5.2 Retrieval Evaluation

- 5.3 BLEU Score

- 5.4 ROUGE Score

- 5.5 Semantic Similarity

- 5.6 Hallucination Rate

- 5.7 Latency Evaluation

- 5.8 Human Evaluation and Error Analysis

- 5.9 Dashboard Metrics and Continuous Testing

- **Chapter 6: Deployment, Security, and MLOps**

- 6.1 Deployment Strategy: VM Trial and Docker Containerization

- 6.2 Azure Container Registry and Web App Deployment

- 6.3 API Authentication and Security

- 6.4 Monitoring Dashboard

- 6.5 MLflow Experiment Tracking

- 6.6 User Feedback Monitoring

- 6.7 Retraining Pipeline

- 6.8 Continuous Improvement Workflow

- **Chapter 7: Business KPI Impact Analysis**

- 7.1 Final System Performance

- 7.2 Response Time Improvement

- 7.3 Ticket Reduction Analysis

- 7.4 Customer Satisfaction Impact

- 7.5 Support Workload Reduction

- 7.6 Cost Reduction Analysis

- 7.7 Human Support vs AI-Assisted Support

- 7.8 Overall Business Impact

- **Chapter 8: Conclusion and Future Work**

- 8.1 Conclusion

- 8.2 Challenges Faced

- 8.3 Lessons Learned

- 8.4 Future Improvements

Abstract

This project presents a Retrieval-Augmented Generation (RAG)-powered intelligent chatbot developed for CityFoam to automate customer support and provide accurate, real-time responses to customer inquiries. The system combines Azure OpenAI with the multilingual sentence-transformers/paraphrase-multilingual-MiniLM-L12-v2 embedding model and ChromaDB vector storage to retrieve relevant knowledge-base content before generating grounded answers. The chatbot supports Arabic and English queries, maintains short-term conversational context, and restricts its answers to retrieved company knowledge in order to reduce hallucination risk and improve response reliability.

The solution includes a FastAPI backend, a web-based chat interface, an administrative dashboard, document ingestion and preprocessing pipelines, Docker-based containerization, Azure Container Registry, and Azure Web App deployment. During deployment planning, a Virtual Machine approach was initially explored; however, Docker containerization was adopted for the final deployment because it provided better portability, reproducibility, dependency management, and cloud-readiness. The system also integrates monitoring and evaluation components, including BLEU, ROUGE, semantic similarity, retrieval hit rate, latency tracking, hallucination detection, user feedback logging, MLflow experiment tracking, and PSO-based retrieval optimization.


Experimental evaluation showed strong performance, including a retrieval hit rate of 93%, semantic similarity of 73%, average latency of 2.7 seconds, and hallucination rate of 0.83%. The KPI dashboard further demonstrated the business value of the chatbot through faster response time, 24/7 availability, reduced support workload, and lower estimated operational cost compared with traditional human-only support. Overall, the project demonstrates how a RAG-based AI support system can combine technical reliability with measurable business impact, while still preserving human support for complex or sensitive cases.

Chapter 1: Introduction

1.1 Project Overview:

CityFoam requires a scalable customer support solution capable of answering product, policy, branch, and service-related inquiries in real time. This project develops a RAG-powered intelligent chatbot that combines semantic retrieval with large language model generation to provide accurate, context-aware, and multilingual responses.

The system uses ChromaDB as a vector database, paraphrase-multilingual-MiniLM-L12-v2 for embedding generation, Azure OpenAI for response generation, and FastAPI for serving the chatbot through REST APIs. A web interface and admin dashboard support user interaction, monitoring, retraining, and performance evaluation. The application is containerized using Docker and deployed on Azure Web App Service to support production-readiness and scalability.



1.2 Problem Statement:

Traditional customer support systems depend heavily on human agents, which limits availability, increases operational cost, and makes scaling difficult during high-demand periods. Rule-based chatbots partially reduce workload but often fail to understand user intent, support multilingual queries, or provide answers grounded in updated company knowledge.

For CityFoam, customers may ask about products, prices, branches, refund policies, and delivery-related information. Manually handling these inquiries can be slow and inconsistent. Therefore, the project aims to provide an intelligent support system that retrieves relevant information from official company documents and generates reliable answers in real time.

Support Aspect	Human-based support	AI-driven automated support
Availability	Limited working hours	24/7 availability
Response Time	Slow response time	Real-time responses
Operational Cost	High operational cost	Reduced operational cost
Scalability	Difficult to scale	Highly scalable
Information Retrieval	Manual information search	Automated semantic retrieval
Context Understanding	Depends on agent experience	Context-aware AI understanding
Response Consistency	Inconsistent responses	Consistent responses
Multilingual Capability	Limited multilingual support	Multilingual support
Maintenance Requirement	Requires continuous staff training	Requires periodic model updates
Concurrent Request Handling	Limited concurrent handling	Handles multiple users simultaneously
Error Rate	Higher human error rate	Lower hallucination and error rate
Customer Satisfaction	Long customer waiting times	Faster customer satisfaction
Employee Workload	Heavy workload on employees	Reduced workload on support teams
System Intelligence Level	Static rule-based chatbots	Dynamic RAG-powered generation

1.3 Objectives:

The main objectives of this project are to:

- Build a RAG-based chatbot for CityFoam customer support.
- Retrieve accurate answers from company documents using vector search.
- Support Arabic and English customer queries.
- Generate grounded responses using Azure OpenAI.
- Provide REST API integration through FastAPI.
- Monitor latency, satisfaction, retrieval quality, and hallucination risk.
- Support retraining and continuous improvement.
- Deploy the system using Docker and Azure Web App Service.

1.4 Scope of the Project:

This project focuses on the design, development, deployment, and evaluation of a Retrieval-Augmented Generation (RAG)-powered intelligent chatbot for CityFoam customer support. The system is intended to assist customers by answering questions related to company products, prices, branches, refund policies, delivery information, and general support services using information retrieved from the company's internal knowledge base.

The scope of the project covers the full lifecycle of an AI-powered customer support system, starting from data collection and preprocessing, moving through semantic retrieval and response generation, and ending with monitoring, retraining, deployment, and business impact analysis.

The project scope includes the following main components:

1.4.1 Data Collection and Knowledge Base Preparation:

The system collects and processes different types of company support data, including product documents, FAQs, support-related files, and structured product information. These files are prepared to become the main knowledge source used by the chatbot during response generation.

The preprocessing stage includes:

- Extracting text from multiple file formats.
- Cleaning noisy or inconsistent text.
- Normalizing Arabic text to improve retrieval accuracy.
- Splitting long documents into smaller chunks.
- Preparing the final processed corpus for vector storage.

This stage is important because the quality of the chatbot's answers depends directly on the quality, structure, and clarity of the retrieved knowledge.

1.4.2 RAG-Based Chatbot Development:

The core of the project is the development of a RAG pipeline that combines document retrieval with large language model generation. Instead of relying only on the general knowledge of an LLM, the chatbot first retrieves relevant information from CityFoam's knowledge base and then uses the retrieved context to generate an answer.

The RAG pipeline includes:

- Query preprocessing and optimization.
- Multilingual embedding generation.
- Semantic similarity search using a vector database.
- Retrieval of the most relevant document chunks.
- Context-grounded answer generation using Azure OpenAI.
- Arabic and English response support.
- Short-term conversational memory for follow-up questions.

This approach helps reduce hallucination because the chatbot is instructed to answer based on retrieved company information rather than unsupported assumptions.

1.4.3 System Architecture and API Integration:

The project includes the implementation of a backend service that exposes the chatbot through REST APIs. FastAPI is used to manage user requests, validate inputs, run the retrieval pipeline, generate responses, and return answers to the frontend interface.

The system architecture includes:

- A FastAPI backend service.
- A ChromaDB vector database.
- An Azure OpenAI integration layer.
- A document ingestion pipeline.
- A web-based chatbot interface.
- An administrative dashboard for monitoring and retraining.
- API-based communication between system components.

This makes the chatbot suitable for integration with a customer support portal or business-facing web application.

1.4.4 Monitoring, Evaluation, and Continuous Improvement:

The project includes monitoring and evaluation mechanisms to measure both technical performance and user-facing quality. The chatbot is evaluated using retrieval, generation, latency, and satisfaction-related metrics.

The monitoring and evaluation scope includes:

- BLEU score.
- ROUGE score.
- Semantic similarity.
- Retrieval hit rate.
- Response latency.
- Hallucination rate.
- User satisfaction feedback.
- Monitoring dashboard.

- MLflow-based experiment tracking.
- Comparison between model performance before and after retraining.

In addition, the project includes a retraining mechanism that allows the system to refresh embeddings and improve retrieval quality when the knowledge base is updated or when performance degradation is detected.

1.4.5 Retrieval Optimization Using PSO:

The scope also includes Particle Swarm Optimization (PSO) as an advanced optimization technique for improving the retrieval and ranking behavior of the RAG system. PSO is used to tune important retrieval and ranking parameters instead of relying only on fixed manually selected values.

The optimized parameters include:

- Number of retrieved documents.
- Similarity threshold.
- Chunk weighting.
- Semantic score weight.
- Keyword score weight.
- Recency score weight.
- Authority score weight.

This allows the system to improve its retrieval quality based on real or evaluation-based query patterns.

1.4.6 Deployment and Security:

The project includes containerization and cloud deployment to support production readiness. Docker is used to package the application and its dependencies, while Azure Web App Service is used for hosting.

The deployment and security scope includes:

- Docker-based containerization.
- Azure Web App Service deployment.

- Environment-based configuration.
- API key authentication.
- Secure management of sensitive credentials.
- Backend validation using Pydantic models.
- Controlled access to system endpoints.

These components help make the system more scalable, maintainable, and suitable for real-world usage.

1.4.7 Business KPI Impact Analysis:

The project also evaluates the expected business value of deploying the chatbot in a customer support environment. The analysis focuses on how automation can improve customer service efficiency and reduce operational pressure on human agents.

The KPI analysis includes:

- Response time improvement.
- Customer satisfaction impact.
- Ticket volume reduction.
- Support workload reduction.
- Operational cost reduction.
- Availability improvement through 24/7 support.
- Comparison between human-based support and AI-assisted support.

This part connects the technical implementation with measurable business value.

Project Boundaries:

Although the system is designed to support real customer support scenarios, its scope is limited to answering questions based on the available CityFoam knowledge base. The chatbot does not independently invent product information, process real payments, modify customer orders, or replace all human support operations. Instead, it acts as an intelligent first-line support assistant that can answer common inquiries and reduce the workload on support teams.

Complex cases, missing information, complaints requiring human judgment, or requests outside the available knowledge base should still be escalated to human agents.

1.5 Technologies Used:

The development of the proposed chatbot system required the integration of multiple technologies and frameworks to ensure high performance, scalability, and production readiness. The following technologies were utilized throughout the project:

Technology	Purpose
FastAPI	Backend API development and RAG pipeline orchestration
Azure OpenAI	Query optimization and grounded response generation
SentenceTransformers MiniLM	Multilingual embedding generation for Arabic and English queries
ChromaDB	Persistent vector database for semantic search and document retrieval
PyMuPDF4LLM	PDF text extraction and conversion into LLM-friendly Markdown
Unstructured	DOCX and TXT document parsing for knowledge base preparation
Pandas	Excel and CSV data extraction for structured product information
HTML, CSS, JavaScript	Web-based chatbot interface and administrative dashboard
Docker	Application containerization and dependency packaging
Azure Web App Service	Cloud hosting and production deployment
MLflow	Experiment tracking, monitoring, and performance comparison
Particle Swarm Optimization (PSO)	Optimization of retrieval and ranking parameters
Pydantic	Request validation and API data modeling
Python	Core backend development and system integration

The system also integrates evaluation and monitoring mechanisms for tracking chatbot performance, including BLEU score, ROUGE score, semantic similarity, latency, retrieval hit rate, and hallucination rate. These technologies collectively contribute to building a scalable, intelligent, and production-ready customer support solution.

Chapter 2: Data Collection and Preprocessing

2.1 Data Sources:

The data collection stage represents the foundation of the CityFoam RAG-powered chatbot system. Since the chatbot is designed to answer customer support questions based on company knowledge, the quality and structure of the input data directly affect the accuracy of the generated answers.

In the implemented CityFoam system, the initial knowledge base was built from representative company documents covering the most important customer support areas: product information, refund policy, and branch information.

The implemented data sources are:

- **Product and price catalog (word sheet):**
Used to answer product-related and pricing questions, such as mattress prices, product categories, available sizes, and product specifications.
- **Refund policy document (pdf):**
Used to answer policy-related questions, such as refund conditions, exchange rules, warranty-related inquiries, and cases requiring a receipt or product inspection.
- **Branches spreadsheet (excel):**
Used to answer location-related questions, such as branch addresses, working hours, and nearest branch inquiries.

These data sources collected through automated text scraping from the Ricrac and Wonderland websites. This scraped content was used to enrich the knowledge base with additional product descriptions and feature-related information, supporting broader product inquiry coverage.

These sources were selected because they represent the most common types of customer support questions expected in a retail business such as CityFoam:

- Pricing inquiries.
- Product feature questions.
- Refund and exchange policy questions.
- Warranty-related questions.
- Branch and location questions.

The knowledge base files used in the project are stored inside the project **"/data"** directory. This mixture of PDF, Word, and Excel files made the preprocessing pipeline more realistic because the system needed to handle both structured and unstructured data formats.

2.2 Knowledge Base Preparation:

Before the chatbot can retrieve accurate answers, the raw company files must be converted into a searchable knowledge base. This preparation stage transforms business documents into clean text units that can later be embedded and stored in the vector database.

The knowledge base preparation process follows four main goals:

1. **Convert raw files into machine-readable text.**

The chatbot cannot directly search inside PDFs, Word files, or spreadsheets unless their contents are extracted into plain text.

2. **Preserve the meaning of business data.**

Product names, prices, policy rules, branch addresses, and working hours must remain connected to their original context.

3. **Prepare the text for semantic retrieval.**

The text must be normalized and divided into smaller chunks so the embedding model can represent each part effectively.

4. **Attach metadata to every processed unit.**

Metadata such as file name, document type, and chunk index helps the system trace each retrieved answer back to its source.

The project uses an offline ingestion pipeline implemented in **"src/ingest.py."** This script is responsible for reading files from the configured data folder, processing them according to file type, and building the vector database used by the runtime RAG pipeline.

A key design choice in this stage is separating **offline ingestion** from **online question answering**. The ingestion pipeline runs when the knowledge base is created or updated, while the chatbot API uses the already-built ChromaDB collection during user interaction. This reduces response latency because expensive preprocessing and embedding generation are not repeated for every user query.

2.3 Data Ingestion Pipeline:

The ingestion pipeline is responsible for scanning the data folder, detecting file types, extracting text, normalizing Arabic content, splitting long documents into chunks, generating embeddings, and storing the final results in ChromaDB.

The implementation supports multiple file formats through three processing routes:

Route 1: Spreadsheets:

Spreadsheet files such as .xlsx, .xls, and .csv are processed using **Pandas**. Each row is transformed into a readable text record using a key-value structure. This approach preserves structured information such as branch names, addresses, working hours, and contact details.

Example output format:

```
--- Database Record (zbranches.xlsx) ---  
Branch: ...  
Address: ...  
Working Hours: ...  
Phone: ...
```

This format is useful because it makes tabular data easier for both the embedding model and the LLM to interpret.

Route 2: PDF Documents

PDF files are processed using **PyMuPDF4LLM**, which converts the PDF into Markdown-style text. This is especially important for `productsandprices.pdf`, because product catalogs often contain tables, visual layouts, and price columns.

Using Markdown conversion helps preserve relationships between product names, sizes, and prices better than simple plain-text extraction. This reduces the risk of losing critical product information during preprocessing.

Route 3: Word and Text Documents

Word and text files such as .docx and .txt are processed using the **Unstructured** library. This allows the system to extract paragraphs, headings, and policy rules from documents such as `refundpolicy.docx`.

The extracted content is then combined, normalized, and split into smaller chunks.

Ingestion Flow

The overall ingestion flow can be summarized as follows:

1. Scan the data/ directory.
2. Detect the file extension.
3. Route the file to the correct parser.
4. Extract readable text.
5. Normalize Arabic characters.
6. Split long text into chunks.
7. Generate vector embeddings.
8. Store documents, embeddings, metadata, and IDs in ChromaDB.

Routes code from "src/ingest.py":

The following excerpt shows how the pipeline scans the data directory and applies a text splitter before routing files based on their type:

```
def process_data_folder():  
    """Scans the data folder and extracts content based on file type."""  
    logger.info(f"Scanning folder: {Config.DATA_DIR}")  
    final_documents = []  
    # We increase chunk size slightly for Markdown tables so they don't get cut in half  
    text_splitter = RecursiveCharacterTextSplitter(chunk_size=1200, chunk_overlap=200)  
  
    if not os.path.exists(Config.DATA_DIR):  
        logger.error(f"Data folder '{Config.DATA_DIR}' does not exist.")  
        return []
```

The spreadsheet route is implemented using Pandas:

```
if ext in ['xlsx', 'xls', 'csv']:
    try:
        df = pd.read_csv(file_path) if ext == 'csv' else pd.read_excel(file_path)
        for _, row in df.fillna("N/A").iterrows():
            row_text = f"--- Database Record ({filename}) ---\n"
            row_text += "\n".join([f"{normalize_arabic(str(k))}: {normalize_arabic(str(v))}" for k, v in
row.items()])
            final_documents.append({"text": row_text.strip(), "metadata": {"source": filename,
"type": "spreadsheet"}})
    except Exception as e:
        logger.error(f"Error processing spreadsheet {filename}: {e}")
```

The PDF route uses PyMuPDF4LLM to convert the catalog into Markdown before chunking:

```
elif ext == 'pdf':
    try:
        md_text = pymupdf4llm.to_markdown(file_path)
        clean_text = normalize_arabic(md_text)
        chunks = text_splitter.split_text(clean_text)
        for i, chunk in enumerate(chunks):
            final_documents.append({"text": chunk, "metadata": {"source": filename, "type":
"pdf_as_markdown", "chunk_index": i}})
    except Exception as e:
        logger.error(f"Error processing PDF {filename}: {e}")
```

The Word/text route uses Unstructured:

```
elif ext in ['docx', 'txt']:
    try:
        elements = partition(filename=file_path)
        full_text = "\n\n".join([normalize_arabic(el.text) for el in elements if hasattr(el, 'text')])
        chunks = text_splitter.split_text(full_text)
        for i, chunk in enumerate(chunks):
            final_documents.append({"text": chunk, "metadata": {"source": filename, "type":
"document", "chunk_index": i}})
    except Exception as e:
        logger.error(f"Error processing document {filename}: {e}")
```

2.4 Text Cleaning and Arabic Normalization:

Arabic text requires special preprocessing because the same word may appear in several visually or linguistically similar forms. For example, different Alef forms such as **أ**, **إ**, and **آ** may refer to the same underlying character but can produce different embeddings if treated as separate symbols.

To reduce spelling mismatches and improve semantic retrieval, the system applies Arabic normalization before storing documents in the vector database.

The normalization function performs the following operations:

- Converts Alef variants **أ**, **إ**, and **آ** into plain **ا**.
- Converts Ta Marbuta **ة** into Ha **هـ**.
- Converts Alef Maqsura **ى** into Ya **ي**.
- Removes Arabic diacritics such as Fatha **َ**, Damma **ُ**, Kasra **ِ**, and Tanween **ً**.

This is particularly important for CityFoam because many customer questions are expected to be written in informal Arabic, with inconsistent spelling, missing diacritics, or mixed Arabic-English input.

Arabic Normalize code from "src/ingest.py":

```
def normalize_arabic(text):  
    """Standardizes Arabic characters to eliminate common spelling mismatches."""  
    if not text: return ""  
    text = re.sub(r'['\u064B-\u065F]', '', text)  
    text = re.sub(r'س', 'س', text)  
    text = re.sub(r'ي', 'ي', text)  
    text = re.sub(r'['\u064B-\u065F]', '', text)  
    return text
```

2.5 Tokenization and Chunking:

After text extraction and normalization, long documents must be divided into smaller units before embedding. This is necessary because RAG systems do not usually retrieve entire documents. Instead, they retrieve the most relevant chunks and pass them to the LLM as context.

The CityFoam system uses RecursiveCharacterTextSplitter with the following settings:

```
text_splitter = RecursiveCharacterTextSplitter(chunk_size=1200, chunk_overlap=200)
```

The chunking configuration was selected for two main reasons:

1. Preserving enough context inside each chunk:
A chunk size of 1200 characters allows product details, policy rules, and branch information to remain meaningful rather than being broken into very small fragments.
2. Reducing context loss between chunks:
A chunk overlap of 200 characters allows nearby information to appear in adjacent chunks. This helps preserve continuity when an important sentence is located near a chunk boundary.

The project uses character-level recursive splitting rather than simple fixed-length cutting. Recursive splitting attempts to separate text at natural boundaries such as paragraphs, line breaks, and spaces. This produces more coherent chunks and improves retrieval quality.

For example, instead of cutting a refund policy sentence in the middle, the splitter attempts to preserve complete paragraphs or meaningful sections when possible.

2.6 Exploratory Data Analysis Summary:

Exploratory Data Analysis (EDA) was performed to validate the readiness of the processed knowledge base before connecting it to the final RAG response generation pipeline. The goal of this stage was not only to inspect the available documents, but also to test whether the vector database could retrieve the correct source file for realistic customer support questions.

The EDA process was implemented in the project deliverables notebook:

"deliverables/eda_analysis.ipynb"

The notebook includes corpus profiling, query distribution analysis, golden query dataset creation, retrieval testing, error analysis, and vector-space visualization.

The EDA focused on four main questions:

- How many processed chunks were stored in the vector database?
- Are the chunks distributed across the expected knowledge sources?
- Do the test queries represent realistic customer support intents?
- Can ChromaDB retrieve the correct source document for each query?


2.6.1 Corpus Profiling:

The first EDA step connected to the ChromaDB collection and loaded the stored chunks with their metadata. This helped verify that the ingestion pipeline successfully converted the raw documents into retrievable text chunks.

The ingestion pipeline is designed to read raw files, extract text, normalize Arabic content, split the text into chunks, generate embeddings, and store the final output in ChromaDB.

Connecting to ChromaDB...

Successfully loaded 13 chunks from the database!



This figure shows two important aspects of the processed corpus. The first chart shows how the stored chunks are distributed across the source files, such as the product catalog, refund policy document, and branches spreadsheet. The second chart shows the distribution of chunk lengths in characters.

This analysis is important because extremely short chunks may lack enough context, while extremely long chunks may reduce retrieval precision. The observed chunk distribution confirms that the corpus was split into manageable text units suitable for semantic search.

2.6.2 Golden Query Dataset:

A golden query dataset was created to test the retrieval behavior of the system using realistic customer questions. The dataset contains **21 representative queries** covering English, Arabic, and Franco-Arabic input styles.

The purpose of this dataset is to simulate how real customers may ask questions in different forms, rather than only testing clean or formal queries.


Golden Query Dataset created successfully with 21 queries!

The dataset includes five main customer support categories:

Category	Number	Purpose
Pricing	5	Testing product price and size inquiries
Policy	5	Testing refund, exchange, and return-policy questions
Location	4	Testing branch location and working-hour questions
Product Features	4	Testing product type, material, and feature questions
Warranty	3	Testing warranty duration and coverage questions

2.6.3 Common Query Intent Distribution:

The golden query dataset was visualized to understand the balance between the main customer support intents.



This pie chart shows the distribution of the prepared test queries across the five main support categories. Pricing and policy queries represent the largest portions because product prices and return/exchange conditions are usually among the most frequent customer concerns in a retail support chatbot.

Location, product feature, and warranty questions were also included to ensure that the chatbot is not evaluated only on pricing or policy data, but across different support scenarios.

This query distribution supports the project objective of building a chatbot capable of handling multiple types of CityFoam customer inquiries.

2.6.4 Initial Retrieval Testing:

After preparing the golden query dataset, each query was embedded using the same multilingual MiniLM model used in the ingestion pipeline. The query vector was then sent to ChromaDB to retrieve the top relevant chunks.

The retrieval test checked whether the expected source file appeared among the retrieved results. For example, pricing queries were expected to retrieve productsandprices.pdf, policy queries were expected to retrieve refundpolicy.docx, and branch-related queries were expected to retrieve zbranches.xlsx.

```
Connecting to ChromaDB...
Loading embedding model (this might take a few seconds)...
Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.
Loading weights: 100%|██████████| 199/199 [00:00<00:00, 2730.94it/s]

--- Running Retrieval Tests ---
MISS: 'What is the price of the Perla Pocket mattress?'
      Expected: productsandprices.pdf | Got: ['refundpolicy.docx', 'refundpolicy.docx', 'refundpolicy.docx', 'zbranches.xlsx', 'zbranches.xlsx']
MISS: 'How much is a king size bed?'
      Expected: productsandprices.pdf | Got: ['refundpolicy.docx', 'zbranches.xlsx', 'refundpolicy.docx', 'refundpolicy.docx', 'refundpolicy.docx']
HIT: 'يكام مرتبة ميديكال مقاس 120؟'
HIT: 'ما هي أسعار المراتب السويت المنفصلة؟'
HIT: 'bkam martabet perla?'
HIT: 'Can I return my pillow after 10 days?'
HIT: 'Do I need a receipt to get a refund?'
HIT: 'يبلغ أبل المرتبة بد ما فتحها؟'
HIT: 'ما هي سياسة الاسترجاع لو المنتج فيه عيب صناعة؟'
MISS: 'yanfa3 araga3 el martaba ba3d ma fat7taha?'
      Expected: refundpolicy.docx | Got: ['productsandprices.pdf', 'productsandprices.pdf', 'productsandprices.pdf', 'zbranches.xlsx', 'zbranches.xlsx']
HIT: 'Where is the nearest branch in Cairo?'
HIT: 'What are the working hours for the Alexandria store?'
HIT: 'فين أقرب فرع لمبني قوم في مدينة نصر؟'
HIT: 'fen far3 madinet nasr?'
HIT: 'Are there any mattresses with medical foam?'
HIT: 'إيه الفرق بين السويت المنفصلة والمنفصلة؟'
HIT: 'هل متوفر مخدات ميموري قوم؟'
HIT: 'fi marateb tebeya 3andoko?'
HIT: 'How many years is the warranty on the mattresses?'
HIT: 'ضمان مرتبة بيرلا كام سنة؟'
HIT: 'el daman kam sana?'

=====
Final Retrieval Hit Rate: 85.71%
=====
Great score! Your database is ready for the LLM generation phase.
```

The retrieval test achieved an initial hit rate of **85.71%**. This means that the vector database retrieved the correct expected source file for most of the golden test queries.

The test also revealed three missed cases:

- Two English pricing queries were not correctly mapped to the product catalog.
- One Franco-Arabic policy query was incorrectly mapped to product and branch sources.

These misses are useful because they show where the retrieval pipeline can be improved. In particular, pricing-related retrieval can be sensitive when product names or table structures are not represented clearly in the extracted text. Franco-Arabic queries may also require stronger query normalization or query rewriting before retrieval.

This result supports the decision to include later optimization techniques, such as query optimization and PSO-based retrieval tuning, to improve source selection and ranking.

2.6.5 Error Analysis and Pipeline Improvement

The initial EDA results showed that the retrieval pipeline performed well on most policy, location, warranty, and Arabic product-related queries. However, the error analysis revealed a repeated weakness in some quantitative product queries, especially questions asking for exact mattress prices, heights, or product-specific specifications.

By inspecting the raw context chunks retrieved from ChromaDB, the main issue was traced back to the PDF ingestion stage rather than the language model itself. The product catalog contained table-like layouts, but some tables were visually structured using spacing instead of clear grid borders. As a result, the earlier extraction approach did not preserve the relationship between product names, prices, sizes, and specifications accurately.

This caused two major retrieval problems:

- Some product-price information was missing or poorly structured inside the extracted chunks.
- The retriever sometimes returned irrelevant chunks for quantitative product questions, which could lead the generator to produce inaccurate or unrelated answers.


To solve this issue, the PDF ingestion approach was improved by using `pymupdf4llm` to convert the product catalog into Markdown-style text before chunking. Markdown is more suitable for RAG pipelines because it keeps table-like information in a readable textual structure, allowing the vector database and the LLM to better preserve relationships between rows and columns.

This improvement helped the system handle product catalog data more effectively, especially for questions involving prices, heights, sizes, and product specifications. Combined with Arabic text normalization and query optimization, the updated pipeline became more reliable for messy real-world customer queries, including Arabic and Franco-Arabic inputs.

The error analysis therefore showed that retrieval quality depends not only on the embedding model, but also on the quality of document extraction and preprocessing. For this reason, improving the ingestion pipeline was a necessary step before moving to final RAG generation and evaluation.

2.6.6 Vector Space Visualization:

To further inspect the quality of the embeddings, PCA was applied to reduce the high-dimensional document embeddings into a two-dimensional space. This allowed visual inspection of how chunks from different source files were distributed semantically.



This figure visualizes the stored document embeddings after dimensionality reduction using PCA. Each point represents a stored chunk, and each color represents a source file.

The purpose of this visualization is to check whether chunks from different sources occupy distinguishable semantic regions. If chunks from the same source appear close to each other, this suggests that the embedding model is capturing meaningful similarities between related pieces of content.

The visualization supports the suitability of semantic retrieval for this project because the knowledge base sources contain different types of information: product catalog data, policy text, and branch/location records.

2.6.7 EDA Summary:

The EDA stage confirmed that the processed corpus was successfully stored and could be queried through ChromaDB. The knowledge base contained **13 processed chunks**, and the golden query dataset contained **21 representative customer questions** across pricing, policy, location, product features, and warranty categories.

The initial retrieval test achieved a **retrieval hit rate of 85.71%**, showing that the vector database was generally effective in retrieving the correct source document. At the same time, the missed cases revealed useful areas for improvement, especially English pricing queries and Franco-Arabic policy questions.

Overall, the EDA results showed that the knowledge base was ready for the next stage of the RAG pipeline, while also identifying clear optimization opportunities for retrieval accuracy and query handling.

2.7 Final Processed Corpus:

The final processed corpus is the output of the ingestion and preprocessing pipeline. It consists of normalized text chunks, metadata, embeddings, and unique IDs stored inside ChromaDB.

Each processed chunk is stored with four main components:

1. Document text

The cleaned and normalized text extracted from the original source file.

2. Embedding vector

A numerical representation generated using paraphrase-multilingual-MiniLM-L12-v2.

3. Metadata

Information about the source file, document type, and chunk index.

4. Unique ID

A generated identifier used to store and retrieve each chunk from ChromaDB.

The vector database is created using `chromadb.PersistentClient`, which stores the collection persistently on disk. This allows the knowledge base to remain available after application restarts.


```
def build_vector_db(documents):  
    """Embeds and saves documents to ChromaDB."""  
    if not documents:  
        logger.warning("No documents to ingest.")  
        return  
  
    logger.info("Connecting to ChromaDB and embedding documents...")  
    client = chromadb.PersistentClient(path=Config.CHROMA_DB_DIR)  
  
    try:  
        client.delete_collection(name=Config.COLLECTION_NAME)  
    except Exception: pass  
  
    collection = client.create_collection(name=Config.COLLECTION_NAME)  
    embedding_model = SentenceTransformer('paraphrase-multilingual-MiniLM-L12-v2')  
  
    texts = [doc["text"] for doc in documents]  
    embeddings = embedding_model.encode(texts).tolist()
```

The final insertion into ChromaDB is performed as follows:

```
collection.add(  
    documents=texts,  
    embeddings=embeddings,  
    metadatas=[doc["metadata"] for doc in documents],  
    ids=[f"id_{i}" for i in range(len(documents))]  
)  
  
logger.info(f"Successfully stored {len(documents)} chunks in {Config.COLLECTION_NAME}.")
```

This final corpus becomes the semantic memory of the chatbot. During runtime, user questions are embedded and compared against these stored vectors. The most relevant chunks are retrieved and passed to Azure OpenAI as grounded context for answer generation.

By the end of this stage, the system has transformed raw CityFoam business files into a structured, searchable, and production-ready knowledge base suitable for RAG-based customer support.



Chapter 3: System Architecture and RAG Pipeline

3.1 Overall System Architecture:

The CityFoam chatbot system follows a Retrieval-Augmented Generation architecture designed to provide accurate, knowledge-grounded customer support responses. The system combines a FastAPI backend, a persistent ChromaDB vector database, a multilingual embedding model, Azure OpenAI, and a web-based user interface.

The project description requires building a production-ready support automation system using RAG and vector search, serving real queries and integrating with backend systems. It also requires creating a vector store, configuring the RAG pipeline, exposing a REST API, and deploying the chatbot service as part of a support workflow.

The implemented architecture is divided into five main layers:

1. **Frontend Layer**

Provides the web-based chatbot interface used by customers to send questions and receive answers.

2. **API Layer**

Uses FastAPI to expose REST endpoints, validate user requests, call the RAG pipeline, and return chatbot responses.

3. **Retrieval Layer**

Uses ChromaDB to retrieve the most relevant knowledge chunks based on semantic similarity.

4. **Generation Layer**

Uses Azure OpenAI to generate the final response based on the retrieved context.


5. **Monitoring and Administration Layer**

Supports dashboard metrics, user feedback, interaction logging, upload of knowledge files, and retraining workflow.

This chapter focus on describes the system as three tightly connected layers: configuration through "**config.py**", ingestion through "**ingest.py**", and API/RAG execution through "**api.py**". It also states that the system is designed to serve Arabic and English customer questions using Azure OpenAI and a local vector database.

Component	Implementation	Responsibility
Web Chat Interface	HTML, CSS, JavaScript	Allows customers to send questions and receive answers
FastAPI Backend	src/api.py	Receives requests and orchestrates the RAG pipeline
Query Optimizer	Azure OpenAI router prompt	Classifies intent and extracts clean search keywords
Embedding Model	paraphrase-multilingual-MiniLM-L12-v2	Converts queries and chunks into semantic vectors
ChromaDB Vector Store	Persistent ChromaDB collection	Stores and retrieves relevant knowledge chunks
Azure OpenAI	Azure OpenAI deployment	Generates grounded natural-language responses
Monitoring Dashboard	live_metrics.py and dashboard endpoint	Tracks latency, retrieval quality, hallucination risk, and satisfaction
Admin Interface	Admin pages and retraining/upload endpoints	Supports knowledge upload, monitoring, and retraining operations

Full High-Level System Architecture



This diagram shows the full runtime flow of the chatbot. The user interacts with the web interface, which sends the query to the FastAPI backend. The backend normalizes and optimizes the query, converts it into an embedding, retrieves relevant context from ChromaDB, and sends that context to Azure OpenAI to generate the final answer. Monitoring components record latency, retrieval quality, hallucination risk, and user satisfaction.

On the right side, the admin workflow manages the knowledge base and model retraining pipeline. Administrators upload new documents through the admin dashboard, where the data is preprocessed, cleaned, and chunked before generating embeddings and updating ChromaDB. In addition, evaluation both metrics before and after retraining to observe the performance changes.

3.2 Retrieval-Augmented Generation Architecture:


The core intelligence of the system is based on Retrieval-Augmented Generation. Instead of asking the language model to answer directly from its general knowledge, the system first retrieves relevant information from CityFoam's knowledge base, then asks Azure OpenAI to generate an answer using only that retrieved context.

This architecture reduces hallucination because the answer is grounded in the company's own documents. We define RAG as a combination of retrieval, which finds relevant private knowledge, and generation, which uses an LLM to compose a grounded natural-language answer.

The RAG process in the CityFoam system works as follows:

1. The customer sends a question.
2. The backend normalizes the query.
3. The query optimizer identifies the user intent and extracts search keywords.
4. The optimized query is converted into an embedding vector.
5. ChromaDB retrieves the most semantically relevant chunks.
6. The retrieved chunks are combined into a context string.
7. Azure OpenAI generates the final answer based only on that context.
8. The response is returned in the same language as the user query.

This structure separates **knowledge retrieval** from **language generation**, making the chatbot more reliable than a traditional standalone LLM chatbot.



This figure focuses specifically on the RAG workflow. The most important point is that the LLM does not answer in **isolation**. It receives a **carefully constructed prompt** containing the **retrieved knowledge** base context, strict **hallucination** rules, **language** instructions, and conversation **history**.

3.3 Embedding Generation Workflow:

The embedding workflow converts text into numerical vectors so that semantic similarity search can be performed. In this project, both stored document chunks and user queries are represented using the same **multilingual** embedding model:

paraphrase-multilingual-MiniLM-L12-v2

The model is suitable for this system because CityFoam customers may ask questions in Arabic, English, or mixed informal formats. We state that the MiniLM model maps semantically similar sentences into nearby vectors and supports multilingual semantic matching, which is essential for Arabic/English customer support.

The embedding workflow has two phases:

- **Offline Embedding Phase**

This phase happens during data ingestion:

1. Raw files are processed.

2. Text is cleaned and normalized.
3. Long text is split into chunks.
4. Each chunk is encoded into a vector.
5. Vectors are stored in ChromaDB with documents and metadata.

- **Runtime Embedding Phase**

This phase happens when a customer asks a question:

1. The user query is normalized.
2. The query optimizer extracts better search keywords.
3. The optimized query is encoded into a vector.
4. The query vector is compared with stored document vectors.
5. The closest chunks are retrieved.

The use of the same embedding model in both phases ensures that user questions and knowledge base chunks exist in the same semantic vector space.

This excerpt from "**src/api.py**" shows how the optimized user query is converted into an embedding vector before retrieval:

```
query_vector = EMBEDDING_MODEL.encode([optimized_query]).tolist()
```

This vector becomes the input to ChromaDB semantic search to retrieve the closest chunks.

3.4 Vector Store and Semantic Retrieval:

ChromaDB is used as the vector database for storing and retrieving the processed knowledge base. Each stored item contains the text chunk, its embedding vector, metadata, and unique ID.

We can describe ChromaDB as the persistent semantic memory of the system. It stores document embeddings locally and retrieves the most relevant chunks during runtime based on vector similarity.

In the runtime pipeline, the system queries ChromaDB using the query embedding. The number of retrieved chunks depends on the query type:

- **List queries** retrieve more chunks because the user may be asking for multiple items.
- **Branch queries** retrieve enough chunks to cover location-related records.

- **Catalog queries** retrieve more chunks because product catalogs may contain many similar items.
- **General queries** retrieve fewer chunks to keep the response focused.

This adaptive retrieval behavior improves the balance between recall and precision.

```
if is_list_query(user_query): retrieval_limit = 25
elif intent == "branches": retrieval_limit = 7
elif intent == "catalog": retrieval_limit = 25
else: retrieval_limit = 5

results = collection.query(query_embeddings=query_vector, n_results=retrieval_limit)
```

This code shows that retrieval is not fixed for every query. Instead, the system adjusts the number of retrieved chunks according to the detected query intent.

3.5 Query Optimization:

Query optimization is an important stage before semantic retrieval. Real customer questions are often informal, incomplete, misspelled, or written in Arabic, English, or Franco-Arabic. Directly embedding the raw query may retrieve weak or irrelevant chunks.

To reduce this issue, the system uses an Azure OpenAI call as a query **router** and **optimizer**. The optimizer classifies the user query into one of four categories:

- branches
- policy
- catalog
- general

It also extracts focused search keywords from the query. These keywords are then used for embedding and retrieval.

We explain that the query optimizer cleans the query, corrects product-name mismatches, and bridges the gap between user wording and database content.

Why Query Optimization Matters

Query optimization improves retrieval because:

- It removes irrelevant conversational words.
- It identifies the user's intent before search.
- It improves product and policy matching.
- It helps with informal Arabic and Franco-Arabic queries.
- It reduces the chance of retrieving unrelated chunks.

For example, a customer may write:

```
bkam martabet perla?
```

The optimizer can help extract a cleaner search target related to the product **name** and **price** inquiry, allowing the retriever to search more effectively.

The system prompt inside "**src/api.py**" instructs the router to return a strict JSON object with intent and keywords:

```
router_prompt = f"""You are an advanced search router and query optimizer.
Your job is to analyze the user's query and output a strict JSON object with two fields:

1. "intent": Classify the query into EXACTLY one of these categories:
- "branches"
- "policy"
- "catalog"
- "general"

2. "keywords": Clean, focused search keywords extracted from the user's query.

"""
```

The returned intent controls the retrieval behavior, while the extracted keywords are embedded and used to search the vector database.

3.6 Grounded Response Generation:

After retrieving the most relevant chunks from ChromaDB, the system constructs a final prompt for Azure OpenAI. This prompt includes:

- The system identity.
- Strict answer-generation rules.
- The target language.
- Retrieved knowledge base context.
- Previous conversation history.
- The current user question.

The goal is to make the LLM generate an answer that is helpful, concise, and grounded in retrieved CityFoam knowledge.

We state that the final LLM prompt enforces four main rules: zero hallucination, unknown-answer handling, language lock, and context awareness.

Main Generation Rules

1. **Zero Hallucination**

The model must only use facts stated in the retrieved context.

2. **Unknown Answers**

If the answer is not found in the context, the assistant should state that the information is not available.

3. **Language Lock**

The response must be written in the same language as the user query.

4. **Context Awareness**

Previous conversation turns are included to handle follow-up questions.

These rules are essential in customer support because inaccurate answers about prices, branches, refund policies, or warranty conditions can damage customer trust.

The system prompt inside "`src/api.py`":

```
system_prompt = f"""You are the official Customer Support AI Assistant for CityFoam.

Your primary function is to provide accurate, helpful answers based STRICTLY on the official
Knowledge Base provided below.

CRITICAL RULES:

1. ZERO HALLUCINATION: You must only use facts stated in the 'KNOWLEDGE BASE CONTEXT'.
2. UNKNOWN ANSWERS: If the answer cannot be found, politely state that you do not have that
information.
3. LANGUAGE LOCK: You MUST write your ENTIRE reply strictly in {target_language}.
4. CONTEXT AWARENESS: Use the previous conversation history to understand pronouns.
```

This prompt is a central part of the system architecture because it controls how Azure OpenAI uses the retrieved knowledge.

3.7 Conversational Memory:

The chatbot supports short-term conversational memory through the history field in the request body. Instead of storing every conversation permanently inside the backend database, the system expects the frontend to send the previous messages with each new request.

This keeps the API mostly stateless while still allowing the LLM to understand follow-up questions. We describe this as conversation history memory, where the frontend maintains the dialogue turns and resubmits them to the API so the generator can resolve pronouns or implicit references.

For example:

```
User: What is the price of the Perla mattress?

Assistant: ...

User: What about its warranty?
```

In the second question, the word “its” depends on the previous conversation. By including conversation history, the system can understand that the user is still referring to the Perla mattress.

Memory Types in the System

The system uses three forms of memory:

1. **Short-Term Conversation Memory**

Stored in the request payload through the history field.

2. **Long-Term Semantic Memory**

Stored in ChromaDB as document embeddings and retrieved during search.

3. **Parametric Model Memory**

Built into the embedding model and Azure OpenAI model weights, used for language understanding and generation behavior.

This layered memory design allows the chatbot to be both context-aware and grounded in the company knowledge base.

Chat Request Model Code:

```
class ChatMessage(BaseModel):  
    role: str  
    content: str  
  
class ChatRequest(BaseModel):  
    query: str  
    history: list[ChatMessage] = []  
    language: str = "auto"
```

This Pydantic model defines how the API receives the current query and previous conversation turns.

3.8 API Design and Frontend Integration:

The system exposes its functionality through FastAPI endpoints. The main endpoint used by the chatbot is:

Post / api / chat

This endpoint receives the customer question, runs the RAG pipeline, logs monitoring information, and returns the generated response with an interaction ID.

We describe "**api.py**" as the runtime heart of the system because it starts the server, exposes the "**/api/chat**" endpoint, orchestrates the two-stage RAG pipeline, and serves the frontend static UI

The FastAPI backend is responsible for:

- Receiving user queries.
- Validating request data using Pydantic.
- Running query normalization and optimization.
- Generating query embeddings.
- Retrieving relevant chunks from ChromaDB.
- Calling Azure OpenAI for response generation.
- Returning the final answer to the frontend.
- Logging interaction metrics.
- Supporting user feedback through rating.
- Serving dashboard metrics for monitoring.

Endpoint	Purpose	Main Output
GET /	Serves the main chatbot web interface	index.html page
POST /api/chat	Receives user questions and returns RAG-generated answers	Generated answer and interaction ID
POST /api/rate	Records user satisfaction feedback for an interaction	Rating confirmation
GET /api/dashboard/metrics	Returns dashboard metrics for monitoring	Monitoring metrics JSON
POST /api/upload-knowledge	Uploads new knowledge base files	Upload confirmation
POST /api/trigger-retraining	Triggers the retraining and knowledge base refresh workflow	Retraining status and metric comparison
POST /api/admin/login	Authenticates access to the admin interface	Admin authentication token

Chat Endpoint Code:

```
@app.post("/api/chat")
def chat_endpoint(request: ChatRequest):
    start = time.time()

    interaction_id = generate_interaction_id()

    result = get_rag_response(request.query, request.history)

    latency_ms = (time.time() - start) * 1000

    return {
        "status": "success",
        "query": request.query,
        "response": result["answer"],
        "interaction_id": interaction_id
    }
```

This simplified excerpt shows the main logic of the chat endpoint: receive request, call the RAG pipeline, calculate latency, and return the response.

Frontend Integration:

The frontend is implemented using HTML, CSS, and JavaScript. It provides the user interface for customer interaction and communicates with the FastAPI backend through API calls.

The frontend layer supports:

- Sending user messages to /api/chat.
- Displaying chatbot responses.
- Maintaining conversation history.
- Sending satisfaction feedback.
- Accessing the admin dashboard and monitoring views.

The backend also serves static frontend files through FastAPI, allowing the chatbot interface and API to be hosted together in the same application.

```
app.mount("/static", StaticFiles(directory="static"), name="static")

@app.get("/")
def serve_ui():
    return FileResponse("static/index.html")
```

This shows that the backend serves the frontend directly from the "/static" directory.



Chapter 4: Model Development and Optimization

4.1 Embedding Model Selection:

The embedding model is one of the most important components in the CityFoam RAG pipeline because it converts both user queries and document chunks into numerical vector representations. These vectors allow the system to compare the semantic meaning of a customer question with the stored knowledge base, even when the exact words are not identical.

For this project, the selected embedding model is:

sentence-transformers/paraphrase-multilingual-MiniLM-L12-v2

This model was selected because the chatbot is designed to support both Arabic and English customer queries. CityFoam customers may ask questions using formal Arabic, informal Arabic, English, or mixed expressions. Therefore, the embedding model must support multilingual semantic similarity rather than simple keyword matching.

We confirm that this model was chosen because it supports more than one language, maps semantically similar sentences into nearby vector representations, and is lightweight enough for fast inference without requiring a GPU. It is also suitable for question-to-document matching, which is the core retrieval task in the RAG pipeline.

The embedding model is used in two main stages:

1. **During ingestion:**

Each processed document chunk is converted into an embedding vector and stored in ChromaDB.

2. **During runtime:**

Each optimized user query is converted into an embedding vector and compared against the stored document vectors.

Using the same model in both stages ensures that document chunks and user queries are represented in the same semantic vector space.

4.2 Vector Database Configuration:

The vector database stores the processed CityFoam knowledge base after preprocessing and embedding generation. In this project, **ChromaDB** is used as the persistent vector store.

ChromaDB stores each chunk with four main elements:

- The original processed text chunk.
- The embedding vector generated by the MiniLM model.
- Metadata such as source file and document type.
- A unique ID for each stored chunk.

We describe ChromaDB as the long-term semantic memory of the system, where raw knowledge base files are transformed into vector embeddings and saved persistently on disk. During runtime, incoming user questions are compared with these stored vectors to retrieve the most relevant context chunks for answer generation.

The vector database is stored locally using a persistent ChromaDB client. This allows the knowledge base to remain available even after restarting the application.

Code Snip from "**src/api.py**":

```
chroma_client = chromadb.PersistentClient(path=Config.CHROMA_DB_DIR)

try:
    collection = chroma_client.get_collection(name=Config.COLLECTION_NAME)
except ValueError:
    raise RuntimeError(
        f"Collection '{Config.COLLECTION_NAME}' not found. Please run ingest.py first."
    )
```

This code shows that the runtime API loads an existing ChromaDB collection before serving chatbot requests. The collection acts as the searchable knowledge base used by the retriever.

Code Snip from "src/config.py":

```
CHROMA_DB_DIR = os.getenv("CHROMA_DB_DIR", "./chroma_db")  
COLLECTION_NAME = os.getenv("COLLECTION_NAME", "cityfoam_rag")
```

These configuration values make the vector database path and collection name configurable through environment variables.

Stored Element	Description	Role in Retrieval
documents	Processed text chunks extracted from the original knowledge sources	Provides the textual context passed to the LLM after retrieval
embeddings	Numerical vector representations generated by the embedding model	Allows semantic similarity search between queries and stored chunks
metadatas	Source-related information such as file name, file type, and chunk index	Helps trace retrieved chunks back to their original source documents
ids	Unique identifiers assigned to each stored chunk	Allows ChromaDB to manage and retrieve each item reliably

4.3 Retriever Configuration:

The retriever is responsible for selecting the most relevant knowledge chunks from ChromaDB based on the user query. In a RAG system, retrieval quality is critical because the LLM can only generate accurate grounded answers if the retrieved context contains the correct information.

The CityFoam retriever works through the following steps:

1. Normalize the user query.
2. Detect the query language.
3. Use the query optimizer to classify the query intent and extract search keywords.
4. Convert the optimized query into an embedding vector.
5. Query ChromaDB using semantic similarity.
6. Retrieve the most relevant chunks.
7. Combine retrieved chunks into a context string for the LLM.

We are responsible for creating a vector store and configuring the RAG model using embedding-based retrieval or a similar vector search approach.

A key design decision in the implemented retriever is that the number of retrieved chunks is not fixed for all query types. Instead, the system adjusts the retrieval limit according to the query intent as we discussed before in Chapter 3.

4.4 Prompt Engineering Strategy

Prompt engineering controls how the LLM behaves during both query optimization and answer generation. In this project, prompt engineering is used in two main places:

1. **Query Router Prompt**

Used to classify the user query and extract focused search keywords.

2. **Grounded Answer Prompt**

Used to generate the final answer based strictly on the retrieved knowledge base context.

The technical documentation states that the system uses a two-stage RAG pipeline: first, query optimization; second, context-grounded answer generation. It also explains that the final LLM prompt enforces rules such as zero hallucination, unknown-answer handling, language lock, and context awareness.

4.5 Azure OpenAI Integration:

Azure OpenAI is used as the generation and reasoning layer in the CityFoam system. It is used for two main tasks:

1. **Query optimization and intent routing**

The model receives the user query and returns structured intent and keywords.

2. **Final answer generation**

The model receives the retrieved context and produces the final customer-facing answer.

The system uses environment-based configuration for Azure OpenAI credentials and deployment settings. The technical documentation identifies `AZURE_ENDPOINT`, `AZURE_API_KEY`, `AZURE_API_VERSION`, and `AZURE_DEPLOYMENT_NAME` as configuration variables managed through "`config.py`".

Code snip from "`src/api.py`":

```
llm_client = AzureOpenAI(  
    azure_endpoint=Config.AZURE_ENDPOINT,  
    api_key=Config.AZURE_API_KEY,  
    api_version=Config.AZURE_API_VERSION)
```

Code snip from "src/config.py":

```
AZURE_ENDPOINT = os.getenv("AZURE_ENDPOINT")
AZURE_API_KEY = os.getenv("AZURE_API_KEY")
AZURE_API_VERSION = os.getenv("AZURE_API_VERSION", "2024-02-15-preview")
AZURE_DEPLOYMENT_NAME = os.getenv("AZURE_DEPLOYMENT_NAME", "cityfoam-gpt")
```

The system uses different temperature settings depending on the task:

- temperature = 0.0 for query routing because intent classification should be deterministic.
- temperature = 0.2 for final response generation to keep answers natural while still controlled.

4.6 PSO-Based Retrieval Optimization:

The baseline RAG pipeline depends on retrieving relevant chunks from ChromaDB and passing them to Azure OpenAI for grounded answer generation. However, retrieval quality is affected by several numerical parameters, such as how many chunks are retrieved, how strict the similarity threshold is, and how retrieved chunks are ranked before being sent to the LLM.

For this reason, the project uses **Particle Swarm Optimization (PSO)** as an advanced optimization technique to improve retrieval and ranking behavior. Instead of manually testing many parameter combinations, PSO automatically searches for better retrieval settings based on retrieval accuracy and user query patterns. This directly supports the project requirement to optimize RAG components based on relevance and response quality.

We used PSO to optimize several tunable parameters in the CityFoam RAG pipeline, including `top_k`, `similarity_threshold`, `chunk_weight`, `weight_semantic`, `weight_keyword`, `weight_recency`, and `weight_authority`. These parameters control both the initial ChromaDB search and the final ranking of retrieved chunks before generation.

Parameter	What It Controls
top_k	How many candidate documents to retrieve from ChromaDB
similarity_threshold	Minimum semantic relevance score to accept a document
chunk_weight	Influence weight of individual document chunks
weight_semantic	How much the embedding similarity score contributes to the final rank
weight_keyword	How much exact keyword overlap contributes
weight_recency	How much document freshness contributes
weight_authority	How much document authority metadata contributes

4.6.1 Optimization Problem:

The core problem is that there is no single fixed retrieval configuration that works best for all customer queries. A product-price question, a branch-location question, and a refund-policy question may require different retrieval behavior.

For example:

- If top_k is too small, the correct document may not be retrieved.
- If top_k is too large, irrelevant chunks may be added to the context.
- If the similarity threshold is too strict, useful chunks may be filtered out.
- If keyword weight is too high, misspelled or informal queries may fail.
- If semantic weight is too high, exact product names or policy terms may be underweighted.

PSO addresses this issue by treating each possible set of retrieval parameters as a candidate solution. Each candidate is evaluated against query examples, and the optimizer gradually moves toward parameter values that improve retrieval accuracy.

4.6.2 PSO Concept in the CityFoam System:

PSO is based on the idea of a swarm of particles searching for the best solution. In this project, each particle represents one possible configuration of retrieval and ranking parameters.

The PSO documentation maps the swarm concept to the CityFoam system as follows: a particle is one candidate parameter set, the swarm is a group of candidate configurations, the particle position represents values such as top_k and ranking weights, and the quality of a position is measured by retrieval accuracy.

The velocity update equation used in the PSO documentation is:

$$\text{velocity} = (w \times \text{current_velocity}) + (c1 \times r1 \times (\text{personal_best_position} - \text{current_position})) \\ + (c2 \times r2 \times (\text{global_best_position} - \text{current_position}))$$

Where:

- w controls inertia, or how much of the previous movement direction is preserved.
- c1 controls how strongly the particle follows its own best previous solution.
- c2 controls how strongly the particle follows the best solution found by the whole swarm.
- r1 and r2 add randomness to support exploration.

In this project, the PSO cost function is based on retrieval error:

$$\text{cost} = 1 - \text{accuracy}$$

This means that minimizing cost is equivalent to maximizing retrieval accuracy.

4.6.3 Search Parameter Optimization:

The first PSO phase optimizes the parameters that control the initial ChromaDB search. This phase uses a 3-dimensional search space:

- top_k
- similarity_threshold
- chunk_weight

We define this phase as **SearchParameterOptimizer**, which determines how many documents are retrieved, how strict the semantic filtering should be, and how document chunks should influence retrieval quality.

The evaluation process for one candidate parameter set works as follows:

1. Encode the test query into an embedding vector.
2. Query ChromaDB using the candidate top_k.
3. Convert returned distances into semantic scores.
4. Filter out chunks below the candidate similarity threshold.
5. Check whether the expected product or source appears in the retrieved results.
6. Calculate accuracy across all test queries.

The optimizer minimizes:

$$\text{cost} = 1.0 - \text{accuracy}$$

This makes the search phase directly linked to retrieval performance rather than a purely theoretical metric.

4.6.4 Ranking Weight Optimization:

The second PSO phase optimizes how retrieved chunks are ranked after the initial search. This phase uses four ranking weights:

- weight_semantic
- weight_keyword
- weight_recency
- weight_authority

The PSO documentation defines this as **RankingWeightOptimizer**, which searches for the best blend of semantic similarity, keyword overlap, document freshness, and source authority. The ranking score is calculated using the following formula:

$$\text{final_score} = (\text{weight_semantic} \times \text{semantic_score}) + (\text{weight_keyword} \times \text{keyword_score}) \\ + (\text{weight_recency} \times \text{recency_score}) + (\text{weight_authority} \times \text{authority_score})$$

The retrieved documents are then sorted by final_score, and the top-ranked chunks become the context passed to the LLM.

This is useful because customer queries are not always solved by semantic similarity alone. For example, product names, branch names, and policy keywords may require exact or near-exact keyword matching, while other queries may benefit more from semantic similarity.

We also include a penalty mechanism to avoid extreme ranking configurations. For example, if one weight becomes too dominant, the optimizer adds a penalty. This prevents the model from depending only on one signal, such as semantic similarity, while ignoring keyword, recency, or authority signals.

4.6.5 Live Feedback Loop and Parameter Deployment:

The PSO system is designed to improve retrieval using real or collected user interactions. According to the PSO documentation, user queries and extracted expected keywords are logged into `user_queries.jsonl`. The optimizer then reads this interaction data and evaluates different parameter combinations against it.

The feedback loop works as follows:

1. The user sends a query to the chatbot.
2. The RAG pipeline extracts a cleaner keyword or expected target.
3. The interaction is logged into `user_queries.jsonl`.
4. PSO periodically reads the accumulated query data.
5. Candidate parameter sets are evaluated against the logged queries.
6. The best parameters are saved into `pso_best_params.json`.
7. The API loads the updated parameters and uses them for later requests.

We also state that PSO runs as a subprocess rather than inside the request flow. This design prevents optimization from blocking active chatbot requests. If the optimizer crashes, the API can continue serving users because communication happens through the shared configuration file.

PSO Subprocess Trigger Code:

```
def trigger_pso_subprocess():
    script_path = os.path.abspath(__file__)
    subprocess.Popen([sys.executable, script_path, "--child-run"])
```

Scheduler Configuration Code:

```
scheduler.add_job(trigger_pso_subprocess, trigger=IntervalTrigger(hours=2),
    id="pso_hourly_job",replace_existing=True)
```

This design keeps the optimization process separate from the real-time chatbot service.

4.6.6 Zero-Downtime PSO Configuration Update:

A key implementation detail in the PSO design is that optimized parameters are written atomically. This avoids the risk of the API reading an incomplete JSON file while PSO is writing new parameters.

We describe the following atomic update pattern:

```
with open("pso_best_params.tmp.json", "w") as f:
    json.dump(flat_params, f)

os.replace("pso_best_params.tmp.json", "pso_best_params.json")
```

This means the API either reads the old complete configuration or the new complete configuration, but never a partially written file.

The final configuration shape includes both search parameters and ranking weights:

```
"top_k": 8,
"similarity_threshold": 0.3241,
"chunk_weight": 0.7812,
"weight_semantic": 0.6234,
"weight_keyword": 0.2198,
"weight_recency": 0.0941,
"weight_authority": 0.0627
```

The API can load these values on each request, allowing the RAG pipeline to use newly optimized parameters without service restart.

4.6.7 Why PSO Was Used:

PSO was selected because retrieval optimization is not a simple one-variable problem. The quality of retrieved context depends on the interaction between multiple parameters.

We compared PSO with other optimization approaches and explains that manual tuning becomes slow and outdated, grid search becomes computationally expensive, random search does not use previous results effectively, and gradient-based optimization is unsuitable because retrieval accuracy is not differentiable.

For CityFoam, PSO is suitable because:

- It optimizes multiple retrieval and ranking parameters together.
- It uses actual retrieval accuracy as the evaluation signal.
- It can improve over time as user queries accumulate.
- It avoids blocking the API by running in a separate subprocess.
- It supports continuous improvement without manually editing retrieval settings.

At the current stage of the project, PSO has been implemented and tested as a separate optimization module. The experiment successfully produced optimized retrieval and ranking parameters; however, these outputs have not yet been fully connected to the live chatbot API pipeline. For this reason, PSO is documented as a completed experimental component that validates the feasibility of automatic retrieval optimization, while full runtime integration with the production RAG workflow is identified as a future improvement.

4.7 Response Generation Process:

The response generation process is the final stage of the model workflow. At this point, the system has already optimized the query, retrieved relevant knowledge chunks, and prepared the context. Azure OpenAI then generates the final answer using the retrieved context and the system rules.

4.7.1 Final Runtime Flow:

The response generation flow works as follows:

1. The user submits a query through the chatbot interface.
2. The API runs the RAG retrieval process.
3. Relevant chunks are selected from the knowledge base.
4. The LLM receives the retrieved context.
5. The final response is generated.
6. Runtime telemetry is recorded for monitoring and later comparison.

This process connects model development with monitoring because each generated response can be evaluated not only by its content, but also by its latency and user feedback.

4.7.2 Experiment Tracking Link to Response Generation:

Although the final response generation stage uses Azure OpenAI to generate grounded answers from retrieved context, the quality of this response depends strongly on the retrieval configuration used before generation. For this reason, MLflow was used to track RAG experiment variations and compare different retrieval settings before selecting the most suitable configuration.

The tracked experiments focused on three main retrieval parameters: chunk size, chunk overlap, and retrieval K. Each configuration was evaluated using retrieval hit rate, RAG score, poor retrieval percentage, decay score, and top-1 distance. This allowed the team to compare retrieval behavior systematically instead of relying only on manual tuning.

The best observed configuration was chunk900_overlap100_k7, using a chunk size of 900 characters, chunk overlap of 100 characters, and retrieval K equal to 7. This configuration achieved the best hit rate of 80.95% and a best RAG score of 0.3451 in the MLflow experiment dashboard.

Therefore, MLflow is connected to response generation indirectly: it does not replace the RAG pipeline, but it provides evidence for selecting retrieval settings that improve the quality of the context passed to the language model.

Chapter 5: Evaluation and Testing

5.1 Evaluation Methodology:

We evaluated the CityFoam chatbot using a combination of retrieval-based, generation-based, semantic, latency, and reliability metrics. The goal was to measure not only whether the chatbot could generate readable answers, but also whether those answers were grounded in the correct retrieved knowledge.

Since the system is a RAG-based chatbot, evaluation must cover two connected stages:

1. **Retrieval quality**

This checks whether the system retrieves relevant knowledge chunks from the vector database before generation.

2. **Answer quality**

This checks whether the generated response matches the expected answer and remains faithful to the retrieved context.

The evaluation process used a golden test dataset containing representative customer queries. These queries covered the most important CityFoam customer support categories, including pricing, policies, branch locations, product features, and warranty questions.

The evaluation focused on the following metrics:

- Retrieval hit rate.
- BLEU score.
- ROUGE score.
- Semantic similarity.
- Hallucination rate.
- Response latency.
- Human feedback and error analysis.

The evaluation was implemented in "`src/evaluation.py`", while live runtime metrics were logged through "`src/live_metrics.py`". This allowed us to evaluate the system offline using predefined test cases and monitor real user interactions during runtime.

Run Evaluation Code:

```
def run_system_evaluation(rag_chain_wrapper, custom_dataset=None):  
    if custom_dataset is not None:  
        golden_dataset = custom_dataset  
    else:  
        golden_dataset = load_golden_dataset()  
  
    total_bleu = 0  
    total_rouge = 0  
    total_similarity = 0  
    hit_count = 0  
    hallucination_count = 0  
    total_latency = 0
```

This function initializes the evaluation process by loading the golden dataset and preparing accumulators for all metrics. Each test case is passed through the RAG pipeline, then the generated answer is compared with the expected answer.

At the end, the function returns the final metric values:

```
return {  
    "bleu": round((total_bleu / num_cases) * 100, 2),  
    "rouge": round((total_rouge / num_cases) * 100, 2),  
    "hit_rate": round((hit_count / num_cases) * 100, 2),  
    "latency": round(total_latency / num_cases, 3),  
    "semantic_similarity": round((total_similarity / num_cases) * 100, 2),  
    "hallucination_rate": round((hallucination_count / num_cases) * 100, 2)  
}
```

This design gives one unified evaluation output covering generation quality, semantic similarity, retrieval success, latency, and hallucination risk.

Model Evaluation Metrics:

Metric	Result	What it measures	Interpretation
Retrieval Hit Rate	93%	How often the system retrieves or produces an answer supported by the correct knowledge	Strong retrieval performance after preprocessing and optimization
BLEU Score	38%	N-gram overlap between the generated answer and the expected answer	Acceptable score for open-ended customer support answers
ROUGE Score	50%	Longest common sequence similarity between generated and expected answers	Good lexical similarity while allowing different wording
Semantic Similarity	73%	Meaning-level similarity between generated and expected answers	Strong semantic alignment between generated and expected answers
Average Latency	2.7s	Average time required to produce a response	Suitable for near real-time customer support interaction
Hallucination Rate	0.83%	Percentage of responses containing unsupported or suspicious generated information	Low hallucination risk and strong grounding in retrieved context

5.2 Retrieval Evaluation:

Retrieval evaluation measures whether the chatbot retrieves the correct knowledge before generating an answer. In a RAG system, this step is critical because the LLM depends on the retrieved context. If the retrieved chunks are irrelevant, even a strong language model may generate an incomplete or incorrect response.

We evaluated retrieval using representative customer queries and expected sources. The initial EDA retrieval test used 21 golden queries across five categories:

- Pricing.
- Policy.
- Location.
- Product features.
- Warranty.

The initial retrieval test achieved **85.71%** hit rate during the EDA stage. This result showed that the vector database was generally effective, but it also revealed weaknesses in some English pricing queries and Franco-Arabic policy queries.

After improving the ingestion pipeline, applying query optimization, and refining retrieval behavior, the final evaluation reached a **retrieval hit rate of 93%**. This improvement indicates

that the system became more reliable in retrieving the correct knowledge before answer generation.

Retrieval Quality Calculation Code:

```
def compute_retrieval_quality(query: str, chunks: List[str], embed_model) -> float:
    if not chunks:
        return 0.0

    query_emb = embed_model.encode([query])[0]
    chunk_embs = embed_model.encode(chunks)

    dot = np.dot(chunk_embs, query_emb)
    norms = np.linalg.norm(chunk_embs, axis=1) * np.linalg.norm(query_emb) + 1e-8
    similarities = dot / norms

    return round(float(np.max(similarities)) * 100, 2)
```

This function calculates the highest cosine similarity between the user query and the retrieved chunks. A higher score means that at least one retrieved chunk is strongly related to the user query.

5.3 BLEU Score:

BLEU was used to measure lexical overlap between the generated answer and the expected answer. It evaluates how many words or short word sequences in the generated response also appear in the reference answer.

BLEU is useful for checking whether the generated answer uses similar wording to the expected answer. However, it is not enough by itself for evaluating customer support responses, because a correct answer may be phrased differently from the reference answer while still having the same meaning.

We used BLEU as one part of the evaluation, not as the only measure of response quality.

The final BLEU score was: 38%

This result is reasonable for a RAG chatbot because generated answers may be semantically correct even when their exact wording differs from the expected answer.

BLEU Calculation Code:

```
ref_tokens = tokenize_text(eval_expected)
gen_tokens = tokenize_text(eval_generated)

bleu = sentence_bleu(
    [ref_tokens],
    gen_tokens,
    smoothing_function=smoothing.method7
)

total_bleu += bleu
```

The evaluation uses tokenized expected and generated answers. Smoothing is applied to make BLEU more stable for short customer support responses.

5.4 ROUGE Score:

ROUGE was used to measure the overlap between the generated answer and the expected answer from another perspective. We used ROUGE-L, which focuses on the longest common subsequence between both texts.

ROUGE is useful because it captures whether the generated answer preserves the main sequence of important information from the reference answer. Compared with BLEU, ROUGE is often more tolerant when the generated answer is slightly longer or structured differently.

The final ROUGE score was: 50%

This indicates that the chatbot often preserved important answer content while still allowing natural response variation.

ROUGE Calculation Code:

```
scorer = rouge_scorer.RougeScorer(['rougeL'], use_stemmer=True)

rouge_scores = scorer.score(eval_expected, eval_generated)

total_rouge += rouge_scores['rougeL'].fmeasure
```

The evaluation uses ROUGE-L F-measure, which balances precision and recall between the generated answer and the expected answer.

5.5 Semantic Similarity:

Semantic similarity was used to measure whether the generated answer and the expected answer carried the same meaning, even if they used different words. This metric is especially important for a bilingual customer support chatbot because Arabic and English responses may vary in phrasing while still being correct.

We calculated semantic similarity by embedding both the expected answer and the generated answer, then calculating cosine similarity between the two vectors.

The final semantic similarity score was: 73%

This result shows that the chatbot responses were generally aligned with the intended meaning of the expected answers.

Cosine Similarity Calculation Code:

```
def calculate_cosine_similarity(vec1, vec2):

    dot_product = np.dot(vec1, vec2)

    norm_a = np.linalg.norm(vec1)
    norm_b = np.linalg.norm(vec2)

    if norm_a == 0 or norm_b == 0:
        return 0.0

    return float(dot_product / (norm_a * norm_b))
```

The generated and expected answers are embedded first, then compared using cosine similarity:

```
vec_ref = get_embedding(eval_expected)
vec_gen = get_embedding(eval_generated)

sim = calculate_cosine_similarity(vec_ref, vec_gen)
total_similarity += sim
```

Semantic similarity was one of the most important metrics in our evaluation because the chatbot may provide correct answers using different wording from the reference answer.

5.6 Hallucination Rate:

Hallucination evaluation measures whether the chatbot generates information that is not supported by the retrieved context. In customer support, hallucination is a serious issue because incorrect prices, warranty durations, refund rules, or branch details can mislead customers.

We used a strict grounding approach during generation, and we also added a simple hallucination detection function in the monitoring layer. The hallucination detector checks whether the response contains numerical values that do not appear in the retrieved context.

This is useful for CityFoam because many risky hallucinations are numerical, such as:

- Invented prices.
- Wrong warranty periods.
- Incorrect product sizes.
- Unsupported branch details.
- Incorrect delivery or policy durations.

The final hallucination rate was: 0.83%

This low percentage indicates that the chatbot was generally well-grounded in the retrieved knowledge.

Hallucination Calculation Code:

```
def detect_hallucination(answer: str, chunks: List[str]) -> bool:
    if not answer or not chunks:
        return False

    context_text = " ".join(chunks).lower()
    numbers = re.findall(r'\d+', answer)

    for num in numbers:
        if int(num) > 5 and num not in context_text:
            return True

    return False
```

This detector is intentionally conservative and focuses on numerical hallucinations. It does not replace full human evaluation, but it provides a useful automated signal for monitoring response reliability.

5.7 Latency Evaluation:

Latency measures how long the system takes to return a response after receiving a user query. This metric is important because the chatbot is designed for customer support, where users expect fast and interactive responses.

The final average latency was: 2.7 seconds

This result is suitable for real-time customer support use because it allows the chatbot to respond quickly while still performing query processing, retrieval, and LLM generation.

Latency is measured at runtime by recording the start time before running the RAG pipeline and calculating the elapsed time after generation.

Latency Calculation Code:

```
start = time.time()

result = get_rag_response(request.query, request.history)

latency_ms = (time.time() - start) * 1000
```

The latency value is then logged with the rest of the interaction metrics:

```
log_interaction(
    log_path=LOG_FILE,
    interaction_id=interaction_id,
    timestamp=datetime.now(),
    query=request.query,
    response=answer,
    latency_ms=latency_ms,
    error=error,
    intent=intent,
    chunks_retrieved=len(chunks),
    retrieval_quality=retrieval_qual,
    hallucination_flag=hall_flag,
    response_length=response_len,
    prompt_tokens=prompt_tokens,
    completion_tokens=completion_tokens,
    total_tokens=total_tokens
)
```

This allows latency to be monitored continuously rather than measured only once during offline testing.

5.8 Human Evaluation and Error Analysis:

Automated metrics are useful, but they do not fully capture the quality of a customer support chatbot. For this reason, human evaluation and error analysis were also considered when reviewing chatbot responses.

Human evaluation focused on four practical questions:

1. Is the answer factually correct?
2. Is the answer grounded in retrieved CityFoam information?
3. Is the language natural and clear for the customer?
4. Does the response directly answer the user's question?

The system also supports direct user feedback through the **"/api/rate"** endpoint. Users can rate an interaction as positive or negative, and the rating is stored in the live metrics log. This makes it possible to track user satisfaction over time.

```
@app.post("/api/rate")
def rate_interaction(interaction_id: str, rating: int):
    if rating not in (0, 1):
        raise HTTPException(
            status_code=400,
            detail="Rating must be 0 (negative) or 1 (positive)"
        )

    success = update_user_satisfaction(LOG_FILE, interaction_id, rating)

    if not success:
        raise HTTPException(status_code=404, detail="Interaction ID not found")

    return {"status": "success", "message": "Rating recorded"}
```

This endpoint allows the frontend to attach satisfaction feedback to a specific chatbot interaction.

Error Analysis

The main error categories observed during evaluation were:

1. Product Table Extraction Issues:

Some pricing-related errors were caused by product catalog extraction problems. The product catalog contained table-like layouts where prices, sizes, and product names needed to remain connected. If this structure was not preserved correctly, retrieval could return incomplete or irrelevant chunks.

The ingestion pipeline was improved by using Markdown-style PDF extraction, which helped preserve product table relationships more clearly.

2. English Pricing Queries:

Some English price questions were harder to retrieve correctly during the early EDA stage. This happened when English wording did not strongly match the Arabic or structured product information stored in the knowledge base.

Query optimization helped reduce this problem by extracting cleaner search terms before retrieval.

3. Franco-Arabic Ambiguity:

Franco-Arabic queries such as informal transliterated Arabic were more challenging than standard Arabic or English. These queries may contain inconsistent spelling, missing vowels, and mixed-language words.

This issue supports the need for query optimization and future expansion of the test dataset with more real customer examples.

4. Numeric Hallucination Risk:

The most critical hallucination risk was generating unsupported numeric values, such as wrong prices, warranty durations, or product sizes. We reduced this risk by grounding answers in retrieved context and by adding a numeric hallucination detector.

Error Analysis Summary:

Error Type	Cause	Improvement Applied	Expected Impact
Product table extraction issue	PDF tables did not always preserve product-price relationships clearly	Markdown-style PDF extraction was used to improve product table readability	More accurate product-price and specification retrieval
English pricing query mismatch	English wording sometimes differed from the stored product catalog text	Query optimization was added before semantic retrieval	Better handling of English product and pricing questions
Franco-Arabic ambiguity	Informal transliteration created inconsistent query forms	Arabic normalization and query rewriting were used to improve matching	Improved support for informal real-world customer queries
Numeric hallucination risk	Generated numbers may be unsupported if not found in retrieved context	A numeric hallucination detection function was added	Lower risk of unsupported prices, sizes, or warranty values
Incomplete context retrieval	The correct chunk may be missed if retrieval settings are too narrow	Retrieval settings and PSO optimization were used to improve context selection	Higher retrieval coverage and better answer grounding

5.9 Dashboard Metrics and Continuous Testing:

We also used runtime monitoring to evaluate the chatbot after deployment-like interactions. The dashboard metrics are calculated from the live interaction log and summarize the system's operational behavior.

The dashboard tracks:

- Error rate.
- Average latency.
- Average retrieval quality.
- Hallucination rate.
- User satisfaction rate.
- Average response length.
- Total token usage.

The dashboard calculation is implemented in "**get_dashboard_metrics**".

Dashboard Metrics Code:

```
return {  
    "total_interactions": total,  
    "error_rate": round((errors / total) * 100, 2) if total else 0,  
    "avg_latency_ms": round(sum(latencies) / len(latencies), 2) if latencies else 0,  
    "avg_retrieval_quality": round(sum(retrieval_qualities) / len(retrieval_qualities), 2) if  
    retrieval_qualities else 0,  
    "hallucination_rate": round((hallucinations / total) * 100, 2) if total else 0,  
    "user_satisfaction_rate": round((positive / len(rated)) * 100, 2) if rated else None,  
    "total_positive": positive,  
    "total_negative": negative,  
    "avg_response_length": round(sum(response_lengths) / len(response_lengths), 2) if  
    response_lengths else 0,  
    "total_tokens_used": total_tokens,  
    "intents_distribution": intents,  
    "last_updated": datetime.now().isoformat()  
}
```

This makes evaluation continuous rather than limited to offline testing. The same system can be used later to compare different retrieval settings, PSO improvements, retraining results, and user satisfaction trends.

Chapter 6: Deployment, Security, and MLOps

6.1 Deployment Strategy: VM Trial and Docker Containerization:

To prepare the CityFoam chatbot for real-world usage, we focused on deploying the system in a way that supports stability, portability, and maintainability. Before adopting the final container-based deployment strategy, we initially experimented with deploying the application on a **Virtual Machine (VM)**. This early attempt helped us understand the system requirements and the runtime dependencies of the chatbot stack, including the FastAPI backend, the frontend interface, the Python libraries, and the vector database environment.

Although the VM-based approach allowed us to run the application in an isolated environment, it introduced practical limitations. Environment setup was more manual, dependency management was harder to reproduce consistently, and updating the application required more operational effort. In addition, scaling and moving the application between environments was less convenient because the VM approach was tied more closely to the full machine configuration rather than to the application itself.

For these reasons, we moved to **Docker containerization** as the final deployment approach. Docker allowed us to package the application and its dependencies into a portable image that can run consistently across development and production environments. This made the deployment process cleaner and more reproducible.

The Docker-based approach gave us several advantages:

- A consistent runtime environment across local and cloud environments.
- Easier dependency management.
- Simpler packaging of the backend, frontend, and required libraries.
- Better portability and maintainability.
- Faster deployment and update cycles.
- Better integration with Azure container services.

Deployment Aspect	Virtual Machine Approach	Docker Containerization Approach
Environment consistency	More environment-specific and harder to reproduce exactly	Consistent runtime environment across systems
Dependency management	Manual setup and dependency installation	Dependencies are packaged inside the container image
Portability	Less portable because it depends on the full machine image	Highly portable across local and cloud environments
Deployment speed	Slower to configure and update	Faster build, ship, and redeploy cycle
Scalability	Possible but heavier operationally	More suitable for scalable cloud deployment
Maintenance effort	Higher maintenance effort	Lower maintenance effort after containerization
Cloud integration	Less streamlined for modern container workflows	Directly compatible with Azure Container Registry and Azure Web App

6.2 Azure Container Registry and Web App Deployment:


After containerizing the application, we adopted a cloud deployment workflow based on **Azure Container Registry (ACR)** and **Azure Web App**. This approach made it possible to store the application image in a managed container registry and deploy the chatbot as a web application in Azure.

The deployment workflow follows these main steps:

1. Build the Docker image locally.
2. Tag the image for Azure Container Registry.
3. Push the image to ACR.
4. Configure Azure Web App to pull the image from ACR.
5. Start the deployed chatbot service in the cloud.

Using **Azure Container Registry** gave us a centralized place to store and manage the application image. Instead of deploying raw source code directly, we deployed a versioned container image. This improves consistency between testing and deployment, and it makes updating the application more reliable.

Azure Web App was then used to host the running chatbot service. This allowed the frontend and backend to be served through a production-ready managed cloud environment.



Why We Used Azure Container Registry

We used Azure Container Registry because it supports the final container-based deployment model in a structured way. It provides:

- Centralized image storage.
- Version-controlled deployment artifacts.
- Clean separation between build and hosting.
- Better compatibility with Azure Web App deployment.
- Easier update and redeployment flow.

This deployment design also makes future CI/CD automation easier if we later decide to automate image build and release steps.

6.3 API Authentication and Security:

Security is essential for a customer-facing AI chatbot, especially when the system exposes API endpoints and handles runtime requests. In the CityFoam system, security was addressed at multiple levels, including request validation, protected API access, environment-based secret management, and secure handling of suspicious or invalid user inputs.

The security design of the system includes the following elements:

- Environment variables for secrets and configuration.
- Request validation using Pydantic models.
- API key authentication support.
- Controlled access to administrative functionality.
- Protection against unsafe or unsupported user prompts.

- Runtime refusal of requests that attempt to expose hidden information or inject executable code.

6.3.1 Environment-Based Secret Management:

Sensitive values such as Azure credentials and deployment settings are not hardcoded into the application logic. Instead, they are loaded from environment variables. This reduces the risk of exposing credentials in source code and makes deployment configuration easier across environments.

The main environment variables include:

- AZURE_ENDPOINT
- AZURE_API_KEY
- AZURE_API_VERSION
- AZURE_DEPLOYMENT_NAME
- CHROMA_DB_DIR
- COLLECTION_NAME

6.3.2 Request Validation:

The API uses Pydantic models to validate request structure before processing. This ensures that the backend receives properly structured input and reduces the risk of malformed requests causing unexpected behavior.

The request schema includes fields such as the user query, conversation history, and language setting. By validating these inputs before they enter the RAG pipeline, we improve system robustness and reduce avoidable runtime errors.

6.3.3 API Authentication:

The system design includes support for API authentication so that access to protected services can be controlled. This is especially important for administrative or management-related routes such as retraining, file upload, and dashboard access.

Authentication helps protect the system against unauthorized use and prevents unrestricted access to sensitive operational features.

6.3.4 Prompt and Input Safety:

We also validated the chatbot's behavior against unsafe or inappropriate user prompts. The screenshots collected from the application demonstrate that the chatbot refuses to reveal internal knowledge-base secrets and does not execute or process unsafe programming instructions as direct commands.


This behavior is important because it shows that the chatbot remains bounded by its intended support function instead of exposing hidden data or acting on malicious prompt injections.

For examples:

This screenshot shows a user asking the chatbot to send the data rules or internal database rules. The chatbot refuses and clarifies that it cannot share the database or confidential internal information, then redirects the interaction back to supported CityFoam product-related help.



This screenshot shows another attempt to force the chatbot to ignore previous instructions. The resulting interaction demonstrates that the system does not reveal internal logic or hidden configuration information and remains restricted by platform safety behavior.



This screenshot shows a prompt containing a script injection attempt: `<script>alert(1)</script>`. The chatbot refuses to process this kind of input and explains that it cannot handle that type of entry, which confirms that unsafe prompt content does not become executable system behavior.



These screenshots provide practical evidence that the chatbot does not respond to secret extraction attempts or direct script-like injection prompts.

6.4 Monitoring Dashboard:

As we discussed before in Chapter 5, The CityFoam chatbot includes a monitoring dashboard to observe system behavior during operation. The dashboard is connected to the runtime logging layer and summarizes key operational metrics collected from chatbot interactions.

The dashboard helps us track how the system behaves after deployment and provides visibility into both technical performance and user-facing quality.

The dashboard is useful for both technical monitoring and operational decision-making. For example, if latency increases, hallucination rate rises, or user satisfaction drops, these signals can indicate that the system needs investigation, knowledge-base refresh, or retraining.

6.5 MLflow Experiment Tracking:


MLflow was used as part of the MLOps layer to track and compare different RAG retrieval configurations. Instead of relying on a single manually selected retrieval setup, several experiment runs were evaluated using different combinations of chunk size, chunk overlap, and retrieval K. This made it possible to measure how preprocessing and retrieval configuration affected the quality of retrieved context before answer generation.

The MLflow experiment dashboard recorded 18 tested RAG configurations. Each run represented a specific combination of:

- Chunk size.
- Chunk overlap.
- Retrieval K.
- Hit rate.
- RAG score.
- Decay score.
- Poor retrieval percentage.
- Top-1 distance.

The best recorded run was:

- Best run: chunk900_overlap100_k7
- Chunk size: 900
- Chunk overlap: 100
- Retrieval K: 7
- Best hit rate: 80.95%
- Best RAG score: 0.3451
- Lowest poor retrieval percentage: 19.05%



The hit rate represents how often the retriever returned the expected source document for the evaluation queries. The RAG score was used as a retrieval separation indicator, where a better score reflects stronger distinction between relevant and irrelevant retrieval results. Poor retrieval percentage measured weak or failed retrieval cases, where lower values are preferred.

The MLflow dashboard also included visual comparisons across all tested configurations. These visualizations helped identify which settings improved retrieval quality and which settings caused weaker retrieval behavior. For example, the dashboard showed that increasing or decreasing chunk size and retrieval K could significantly affect hit rate and poor retrieval percentage.

Using MLflow in this way helped the project in four main ways:

- It provided structured experiment tracking for RAG configurations.
- It made retrieval performance comparisons visible through charts and tables.
- It supported evidence-based selection of the best retrieval setup.
- It created a reproducible record of tested configurations for future improvement.

This confirms that MLflow was not only used as a logging tool, but also as an experiment management layer for evaluating RAG pipeline quality.

It is important to distinguish between the MLflow RAG experiment hit rate and the final system evaluation hit rate. The MLflow dashboard reports the best retrieval hit rate achieved during retrieval-configuration experiments, which was 80.95%. The final system evaluation later reported 93% after additional improvements, query optimization, and full pipeline evaluation.

ALL RAG EXPERIMENT RUNS										
Run	Chunk Size	Overlap	K	Hit Rate	RAG Score	Decay Score	Poor %	Top-1 Dist	Chunks	Ingest Secs
chunk900_overlap100_k7	900	100	7	80.95%	0.3451	0.2281	23.81%	15.0529	16	2.06
chunk900_overlap200_k7	900	200	7	71.43%	0.3520	0.2322	19.05%	14.9729	17	2.17
chunk600_overlap100_k7	600	100	7	71.43%	0.3076	0.2457	23.81%	14.7965	23	2.95
chunk1200_overlap100_k7	1200	100	7	71.43%	0.3010	0.1941	38.10%	17.0645	13	2.21
chunk1200_overlap200_k7	1200	200	7	66.67%	0.2981	0.1957	33.33%	16.9442	13	1.94
chunk900_overlap200_k5	900	200	5	66.67%	0.2814	0.2322	33.33%	14.9729	17	2.22
chunk900_overlap100_k5	900	100	5	66.67%	0.2806	0.2281	28.57%	15.0529	16	2.20
chunk1200_overlap100_k5	1200	100	5	61.90%	0.2497	0.1941	57.14%	17.0645	13	1.87
chunk1200_overlap200_k5	1200	200	5	61.90%	0.2432	0.1957	57.14%	16.9442	13	2.21
chunk600_overlap100_k5	600	100	5	61.90%	0.2386	0.2457	66.67%	14.7965	23	2.81
chunk600_overlap200_k7	600	200	7	57.14%	0.2881	0.2427	33.33%	14.7384	25	2.40
chunk900_overlap100_k3	900	100	3	57.14%	0.2010	0.2281	66.67%	15.0529	16	2.10
chunk600_overlap200_k5	600	200	5	52.38%	0.2357	0.2427	61.90%	14.7384	25	2.66
chunk1200_overlap200_k3	1200	200	3	52.38%	0.1417	0.1957	80.95%	16.9442	13	1.91
chunk1200_overlap100_k3	1200	100	3	52.38%	0.1359	0.1941	85.71%	17.0645	13	2.01
chunk900_overlap200_k3	900	200	3	47.62%	0.1950	0.2322	80.95%	14.9729	17	2.02
chunk600_overlap200_k3	600	200	3	47.62%	0.1559	0.2427	90.48%	14.7384	25	2.50
chunk600_overlap100_k3	600	100	3	47.62%	0.1542	0.2457	90.48%	14.7965	23	3.30

6.6 User Feedback Monitoring:

User feedback is an important part of the monitoring loop because automated metrics alone do not fully reflect customer experience. A response may be technically grounded but still not satisfy the user if it is unclear, incomplete, or not directly useful.

To address this, the CityFoam chatbot includes a user feedback mechanism through the rating endpoint. After a chatbot interaction, the user can submit a positive or negative rating linked to the interaction ID.

This gives us a direct satisfaction signal that can be monitored over time and aggregated in the dashboard. It also helps identify patterns such as:

- queries that users often rate negatively,
- intents associated with lower satisfaction,
- periods where quality drops,
- cases that may require retraining or knowledge updates.

User feedback supports the system in three ways:

1. Quality monitoring
It provides an additional real-world measure of answer usefulness.
2. Continuous improvement
It helps identify which interactions should be reviewed.
3. Operational insight
It allows us to compare technical quality with human satisfaction.

6.7 Retraining Pipeline:

The CityFoam system was designed to support knowledge-base updates and retraining-related refresh operations. Since customer support information can change over time, the chatbot must be able to incorporate new or updated knowledge sources without requiring a full redesign.

The retraining pipeline is connected to the administrative workflow and includes the following main steps:

1. Upload or add new knowledge files.

2. Re-run the ingestion and preprocessing pipeline.
3. Rebuild the processed chunks and embeddings.
4. Refresh the ChromaDB collection.
5. Re-evaluate the updated system.
6. Compare the new performance with the previous performance.
7. Decide whether to accept or reject the new retrained version.

In this project, retraining does not only mean changing model weights. It mainly refers to updating the RAG knowledge pipeline so that the chatbot reflects the latest available support content. This includes refreshing document embeddings and evaluating whether the new version improves retrieval and response quality.

Retraining becomes useful in situations such as:

- new product catalogs,
- updated refund or warranty policies,
- new branch information,
- performance degradation,
- user feedback revealing repeated failures.

To ensure system reliability, the retraining pipeline includes an automated validation and rollback mechanism. Before applying any updated knowledge base, the system stores a backup of the currently deployed version, including embeddings, processed chunks, and vector database references. After retraining is completed, the system evaluates the new version using predefined evaluation metrics such as BLEU, ROUGE, semantic similarity, hit rate, hallucination rate, and response latency.

The newly retrained version is then compared against the previous production version. If the evaluation results show a significant performance degradation (15% or more in the main evaluation metrics), the retraining process is automatically rejected and marked as failed. In this case, the system performs an automatic rollback by restoring the previously saved stable version of the knowledge base.

Additionally, the uploaded knowledge file associated with the failed retraining attempt is automatically removed to prevent corrupted or low-quality data from remaining in the system. This mechanism ensures that only validated and performance-safe updates are deployed to the chatbot environment while maintaining system stability and response quality.

CityFoam RAG Admin Panel

Upload New Knowledge File

اختار ملف

CEO Data.xlsx

Upload File

Knowledge Base Refresh

Start Retraining

تم اعتماد التدريب وإضافة الأسئلة الجديدة للملف الذهبي

Evaluation Metrics

Metric	Previous	Current	Change
BLEU	37.97%	36.13%	-1.84%
ROUGE	49.7%	47.86%	-1.84%
Hit Rate	93.33%	93.33%	+0%
Latency	2.729s	3.407s	+0.68s
Semantic Similarity	72.44%	71.96%	-0.48%
Hallucination Rate	0.83%	0.83%	+0%

6.8 Continuous Improvement Workflow:

The final MLOps objective of the CityFoam system is not only to deploy the chatbot once, but to keep improving it over time. Continuous improvement in this project is achieved by combining deployment, monitoring, feedback, experiment tracking, retraining, and retrieval optimization.

The workflow operates as follows:

1. The chatbot is deployed as a containerized application.
2. User interactions are logged continuously.
3. Dashboard metrics and MLflow runs monitor operational behavior.
4. User feedback highlights practical quality issues.
5. Retrieval optimization improves ranking and selection behavior.

6. Knowledge updates trigger retraining or corpus refresh.

7. Updated versions are redeployed through the container workflow.

This creates a closed improvement loop where deployment and operations are connected directly to future system quality.

Chapter 7: Business KPI Impact Analysis

7.1 Final System Performance:

The CityFoam chatbot was evaluated not only from a technical perspective, but also from a business impact perspective. Since the system is designed to automate customer support, the most important business question is whether the chatbot can improve response speed, availability, cost efficiency, and operational workload compared with traditional human-only support.


The KPI dashboard was used to compare the chatbot against estimated human customer support benchmarks. The dashboard includes response time, no-hallucination accuracy, user satisfaction, retrieval quality, availability, token usage, and estimated monthly cost.

The results shown in the dashboard represent a prototype testing snapshot. Some values, especially user satisfaction, were generated during testing using limited or random feedback entries. Therefore, they should be interpreted as experimental dashboard results rather than final production customer satisfaction. In real usage, satisfaction is expected to improve as the chatbot receives more realistic feedback, richer knowledge-base coverage, and further retrieval optimization.

ChatAdmin


CityFoam Performance Dashboard

Live metrics comparing AI chatbot vs Human customer support benchmarks.



This figure presents the main KPI cards and charts used to compare the CityFoam chatbot with human support benchmarks. The dashboard shows that the chatbot achieved an average response latency of **2.7 seconds**, compared with an estimated **30 seconds** for human support. It also shows no-hallucination accuracy, user satisfaction, and retrieval quality indicators.

The most important observation from this dashboard is that the chatbot provides a major response-time improvement while maintaining strong automated response reliability.



This figure shows the full comparison table between the CityFoam chatbot and human support. The table compares average response time, accuracy, user satisfaction, retrieval quality, availability, token usage, and estimated monthly cost. These values help connect the technical chatbot implementation with measurable business outcomes.

7.2 Response Time Improvement:

Response time is one of the strongest business advantages of the chatbot. In customer support, faster responses reduce waiting time, improve user experience, and allow the business to handle more inquiries at the same time.

The dashboard recorded:

- **Bot average response time:** 2.7 seconds
- **Human average response time:** 30.0 seconds

This means the chatbot was approximately: $30.0 / 2.7 = 11.1$ times faster

The percentage reduction in response time is: $((30.0 - 2.7) / 30.0) \times 100 = 91\%$

This result shows that the chatbot can significantly reduce the time required to answer common customer questions. Instead of waiting for an available support agent, customers can receive an immediate answer from the chatbot.

Business Impact

- Faster customer support interaction.
- Lower waiting time.
- Better handling of peak-time inquiries.
- Reduced pressure on human agents.
- Improved first-response experience.

7.3 Ticket Reduction Analysis:

A major benefit of the chatbot is its ability to reduce the number of repetitive tickets handled by human agents. Many CityFoam customer inquiries are routine and can be answered directly from the knowledge base, such as questions about prices, branches, refund policies, warranties, and product features.

The chatbot can reduce ticket load by automatically handling:

- Product price inquiries.
- Branch location questions.
- Working-hour questions.
- Refund and exchange policy questions.

- Warranty questions.
- Product feature and material questions.
- Repeated general support inquiries.

This does not mean that the chatbot replaces human agents completely. Instead, it works as a first-line support assistant. It handles common inquiries instantly and leaves more complex cases for human agents.

Expected Ticket Reduction Impact:

Area	Expected Impact
Repetitive inquiries	Reduced manual handling
Simple product questions	Answered automatically
Policy questions	Retrieved from knowledge base
Branch questions	Answered instantly
Complex complaints	Escalated to human support

By automating common support questions, the company can reduce human workload and allow agents to focus on complex or sensitive cases that require human judgment.

7.4 Customer Satisfaction Impact:

Customer satisfaction depends on several factors, including speed, accuracy, clarity, availability, and trust. The chatbot improves some of these areas strongly, especially response speed and 24/7 availability.

The dashboard showed:

- **Bot satisfaction:** 66.7%
- **Human satisfaction benchmark:** 85%

This lower bot satisfaction value should not be interpreted as a final production result because the dashboard feedback was based on limited testing and random experimental ratings. In a real production environment, satisfaction is expected to improve as:

- more real users interact with the system,
- the knowledge base becomes richer,
- query optimization improves,
- PSO improves retrieval ranking,
- the chatbot receives more feedback-driven updates,
- the frontend experience becomes more polished.

Even with the prototype satisfaction value, the chatbot already provides strong advantages in speed and availability. These advantages can positively affect satisfaction when the system is deployed with realistic customer usage.

Satisfaction Improvement Factors

- Faster answers.
- Always-available support.
- Consistent responses.
- Reduced waiting time.
- Continuous improvement through user feedback.
- Better answers after retraining and knowledge-base updates.

7.5 Support Workload Reduction:

The chatbot reduces support workload by answering routine questions automatically. This allows human support agents to spend more time on complex cases instead of repeatedly answering the same questions.

In a traditional support process, human agents manually search for product details, prices, branch information, or policy rules. In the chatbot system, these tasks are automated through semantic retrieval and grounded generation.

Workload Reduction Areas:

Support Task	Traditional Handling	Chatbot Handling
Product price inquiry	Agent searches catalog manually	Bot retrieves product information
Branch inquiry	Agent checks branch list	Bot retrieves branch details
Refund question	Agent checks policy document	Bot retrieves refund policy
Warranty inquiry	Agent checks product/policy data	Bot answers from knowledge base
Repeated FAQs	Agent repeats same answer	Bot answers instantly

This workload reduction improves operational efficiency because human agents can focus on complaints, special cases, and customer situations that require negotiation or judgment.

7.6 Cost Reduction Analysis:

Cost reduction is one of the clearest business benefits shown in the dashboard. The comparison table estimated:

- **Bot monthly cost:** approximately \$100 server cost
- **Human support cost:** approximately \$3,000 per agent monthly

The estimated cost saving compared with one human support agent is:

$$((3000 - 100) / 3000) \times 100 = 96.7\%$$

This does not mean that the chatbot eliminates the need for all human agents. Instead, it shows that a chatbot can handle a large portion of repetitive inquiries at a much lower operational cost.

Business Interpretation

The chatbot provides high-cost efficiency because one deployed system can serve many users simultaneously, while human support capacity depends on agent availability and working hours. This makes the chatbot especially valuable for high-volume support scenarios.

Cost Comparison Summary:

Metric	CityFoam Bot	Human Support
Estimated monthly cost	~\$100 server cost	~\$3,000 per agent
Cost scalability	Scales with infrastructure	Scales with number of agents
Availability	24/7	Limited working hours
Concurrent requests	Multiple users	Limited by agent capacity

The main value is not only cost reduction, but also better scalability. The chatbot can handle concurrent customer questions without requiring proportional increases in human staff.

7.7 Human Support vs AI-Assisted Support:

The dashboard comparison shows that each support approach has strengths. Human support remains important for complex and sensitive cases, while the chatbot is more efficient for repetitive, knowledge-based inquiries.

The chatbot performed strongly in:

- Response time.
- Availability.

- Cost efficiency.
- Consistency.
- Handling repetitive inquiries.

Human support remains stronger in:

- Emotional judgment.
- Complex negotiation.
- Exceptional cases.
- Customer complaints requiring human empathy.
- Final decision-making in unclear cases.

Therefore, the best business model is not full replacement, but **AI-assisted customer support**. In this model, the chatbot acts as the first layer of support, while human agents handle escalated cases.

7.8 Overall Business Impact:

The CityFoam chatbot provides measurable business value by improving support speed, reducing operational cost, increasing availability, and decreasing repetitive workload on human agents.

The strongest KPI improvements are:

- **Response time:** reduced from 30.0 seconds to 2.7 seconds.
- **Speed improvement:** approximately 11.1× faster.
- **Response-time reduction:** approximately 91%.
- **Availability:** improved from limited human availability to 24/7 support.
- **Cost efficiency:** estimated monthly cost reduction of approximately 96.7% compared with one human support agent.
- **Workload reduction:** repetitive inquiries can be automated.
- **Operational scalability:** more users can be served without proportional staffing increase.

The dashboard results show that the chatbot is especially strong as a first-line support assistant. It can provide instant answers for common product, branch, policy, and warranty questions, while human agents remain responsible for complex, sensitive, or exceptional cases.

Overall, the system supports a hybrid customer support model where AI improves speed and scalability, and human agents focus on high-value interactions. This makes the CityFoam chatbot a practical business solution rather than only a technical AI prototype.

Chapter 8: Conclusion and Future Work

8.1 Conclusion:

This project developed a Retrieval-Augmented Generation chatbot for CityFoam customer support. The system was designed to answer customer questions using retrieved company knowledge rather than relying only on the general knowledge of a language model.

The final system combines document ingestion, Arabic normalization, chunking, multilingual embeddings, ChromaDB vector storage, Azure OpenAI response generation, FastAPI APIs, a web-based chatbot interface, a monitoring dashboard, MLflow-based RAG experiment tracking, retraining workflow, and deployment using Docker and Azure services.

The main achievement of the project is building a complete RAG-based customer support pipeline that can:

- process company documents,
- retrieve relevant knowledge,
- generate grounded answers,
- support Arabic and English queries,
- monitor runtime quality,
- collect user feedback,
- support retraining,
- and connect technical performance with business KPIs.

The final full-system evaluation showed strong retrieval and answer-generation performance, including a final retrieval hit rate of 93%, semantic similarity of 73%, average latency of 2.7 seconds, and hallucination rate of 0.83%. In addition, MLflow experiment tracking was used to compare RAG retrieval configurations, where the best tracked configuration achieved an experiment-level hit rate of 80.95% using chunk size 900, overlap 100, and retrieval K equal to 7.

From a business perspective, the chatbot showed clear value in response-time improvement, support availability, cost reduction, and workload reduction. The system is not intended to replace human agents completely. Instead, it provides a practical AI-assisted support model where the chatbot handles repetitive questions and human agents focus on complex or sensitive cases.

8.2 Challenges Faced:

Several challenges appeared during the development of the CityFoam chatbot.

1. PDF Table Extraction:

One of the main challenges was extracting product catalog information from PDF files. Some product tables depended on visual spacing rather than clear grid borders, which made early extraction attempts less reliable. This affected product-price and product-specification retrieval.

We addressed this challenge by using Markdown-style PDF extraction, which preserved table-like information more clearly and improved the readability of product data for retrieval and generation.

2. Arabic Text Variations:

Arabic text can appear in many forms because of differences in Alef, Ta Marbuta, Alef Maqsura, and diacritics. These variations can reduce retrieval accuracy if the query and document use different spellings.

We addressed this issue using Arabic normalization before embedding and retrieval.

3. Franco-Arabic Queries:

Some customers may write Arabic using Latin characters, such as informal Franco-Arabic. These queries are harder to match directly with Arabic company documents.

We reduced this issue through query optimization, but further improvement is still possible by collecting more real user examples and expanding the evaluation dataset.

4. Balancing Retrieval Depth:

Retrieving too few chunks may miss the correct answer, while retrieving too many chunks may add noise to the LLM context. This required adaptive retrieval behavior and later PSO-based optimization.

5. Deployment Approach:

At the beginning, we experimented with Virtual Machine deployment. Although it provided direct control, it required more manual setup and was less portable. We later moved to Docker containerization and Azure Container Registry because this approach provided a more consistent, maintainable, and cloud-ready deployment workflow.

6. Monitoring and Continuous Evaluation:

A chatbot is not finished after deployment. It requires continuous monitoring, user feedback, and retraining. Building dashboard metrics, MLflow tracking, and retraining workflows helped address this challenge.

8.3 Lessons Learned:

This project provided several important technical and practical lessons.

1. Data Quality Controls RAG Quality:

The quality of a RAG chatbot depends heavily on the quality of its knowledge base. Even a strong LLM cannot generate accurate answers if the retrieved chunks are incomplete, poorly extracted, or incorrectly structured.

2. Preprocessing Is as Important as Model Selection:

Arabic normalization, PDF extraction, chunking, and metadata design had a major impact on retrieval quality. The project showed that preprocessing decisions can be as important as the choice of embedding model or LLM.

3. RAG Reduces Hallucination but Does Not Remove All Risk Automatically:

Grounding the LLM in retrieved context significantly reduces hallucination risk, but retrieval errors or incomplete context can still lead to weak answers. This is why hallucination monitoring and error analysis are necessary.

4. Monitoring Turns the Chatbot into a Maintainable System:

Runtime logging, dashboard metrics, user feedback, and MLflow tracking make the chatbot easier to evaluate and improve over time. Without monitoring, it would be difficult to know whether the chatbot is improving or degrading.

5. Containerization Improves Deployment Reliability:

The comparison between VM deployment and Docker deployment showed that containerization provides a more reproducible and maintainable deployment process. Using Azure Container Registry also made the cloud deployment workflow more structured.

6. Business KPIs Are Essential for AI Projects:

Technical metrics such as BLEU, ROUGE, and retrieval hit rate are important, but business metrics such as response time, cost reduction, availability, and workload reduction show the real value of the system.

8.4 Future Improvements:

Although the CityFoam chatbot achieved strong results, several improvements can make the system more accurate, scalable, and production-ready.

1. Expand the Knowledge Base

Future work should include adding more CityFoam product files, updated price lists, branch information, FAQs, customer service scripts, and real support conversations. A richer knowledge base will improve answer coverage and reduce cases where the chatbot cannot find enough context to answer.

2. Improve Franco-Arabic Handling

The system can be improved by adding more Franco-Arabic examples to the evaluation dataset and using more advanced query rewriting techniques for transliterated Arabic. This would help the chatbot better understand informal customer queries such as mixed Arabic-English or Arabic written using English letters.

3. Fully Integrate PSO into the Live RAG Pipeline

Although the PSO module was implemented and tested successfully, it has not yet been fully integrated into the live production RAG API workflow. The module produced optimized retrieval and ranking parameters, but future work should connect these outputs directly to the chatbot runtime pipeline.

This would allow the system to automatically apply the best values for parameters such as `top_k`, `similarity_threshold`, `weight_semantic`, `weight_keyword`, `weight_recency`, and `weight_authority` during real user interactions.

4. Automate the Retraining Pipeline

The current retraining workflow can be improved by making it more automatic instead of relying mainly on manual triggering. Future versions can continuously collect successful questions, useful chatbot responses, frequent customer inquiries, and newly uploaded files into a structured database.

The system can then periodically perform:

- knowledge base refresh,
- embedding regeneration,
- retrieval re-evaluation,
- answer-quality testing,
- comparison between old and new performance,
- and automatic deployment only if the new version improves results.

This would make retraining more reliable and allow the chatbot to improve continuously based on real customer behavior.

5. Strengthen Human Escalation

Future versions should include an escalation mechanism that transfers unresolved, sensitive, or low-confidence questions to a human support agent. This would create a stronger hybrid support model where the chatbot handles common inquiries and human agents handle complex cases.

6. Improve Satisfaction Measurement

The current satisfaction dashboard values were collected during prototype testing. A production version should collect real customer feedback over a longer period to measure satisfaction more accurately and identify weak areas in the chatbot experience.

7. Add Role-Based Access Control

Admin features such as file upload, retraining, dashboard access, and monitoring reports can be improved with stronger role-based access control. This would allow different permission levels for admins, developers, and business users.

8. Automate CI/CD Deployment

The Docker and Azure Container Registry workflow can be extended into a full CI/CD pipeline. This would allow automatic building, testing, pushing, and redeployment of new container versions after approved updates.

9. Add More Advanced Analytics

The dashboard can be extended with trend analysis, intent-level satisfaction, failed-query clustering, cost-per-interaction tracking, and most-asked-question reports. These

analytics would help the business understand customer needs and improve both chatbot quality and company support operations.

10. Integrate with Business Systems

Future versions can connect the chatbot to CRM, order tracking, inventory systems, and support ticket platforms. This would allow the chatbot to answer more personalized and operational questions, such as order status, product availability, and customer-specific service requests.

11. Improve Security and Prompt-Injection Testing

Future work can include more systematic security testing against prompt injection, sensitive-data extraction attempts, and unsafe input patterns. This would strengthen the chatbot's reliability before wider production deployment.

12. Continuous Knowledge Refresh

The knowledge base can be refreshed automatically when new files are uploaded or when monitoring shows a drop in retrieval quality or satisfaction. This would ensure that the chatbot remains aligned with the latest product, policy, and branch information.

Final Project Statement

The CityFoam RAG-powered chatbot demonstrates how AI can be used to improve customer support operations through grounded generation, semantic retrieval, monitoring, and continuous improvement. The project successfully connects machine learning engineering with practical business value by reducing response time, improving availability, lowering operational cost, and supporting a scalable AI-assisted customer support model.